

Databases at scale

DATA 604

Leanne Wu

lewu@ucalgary.ca

Department of Computer Science

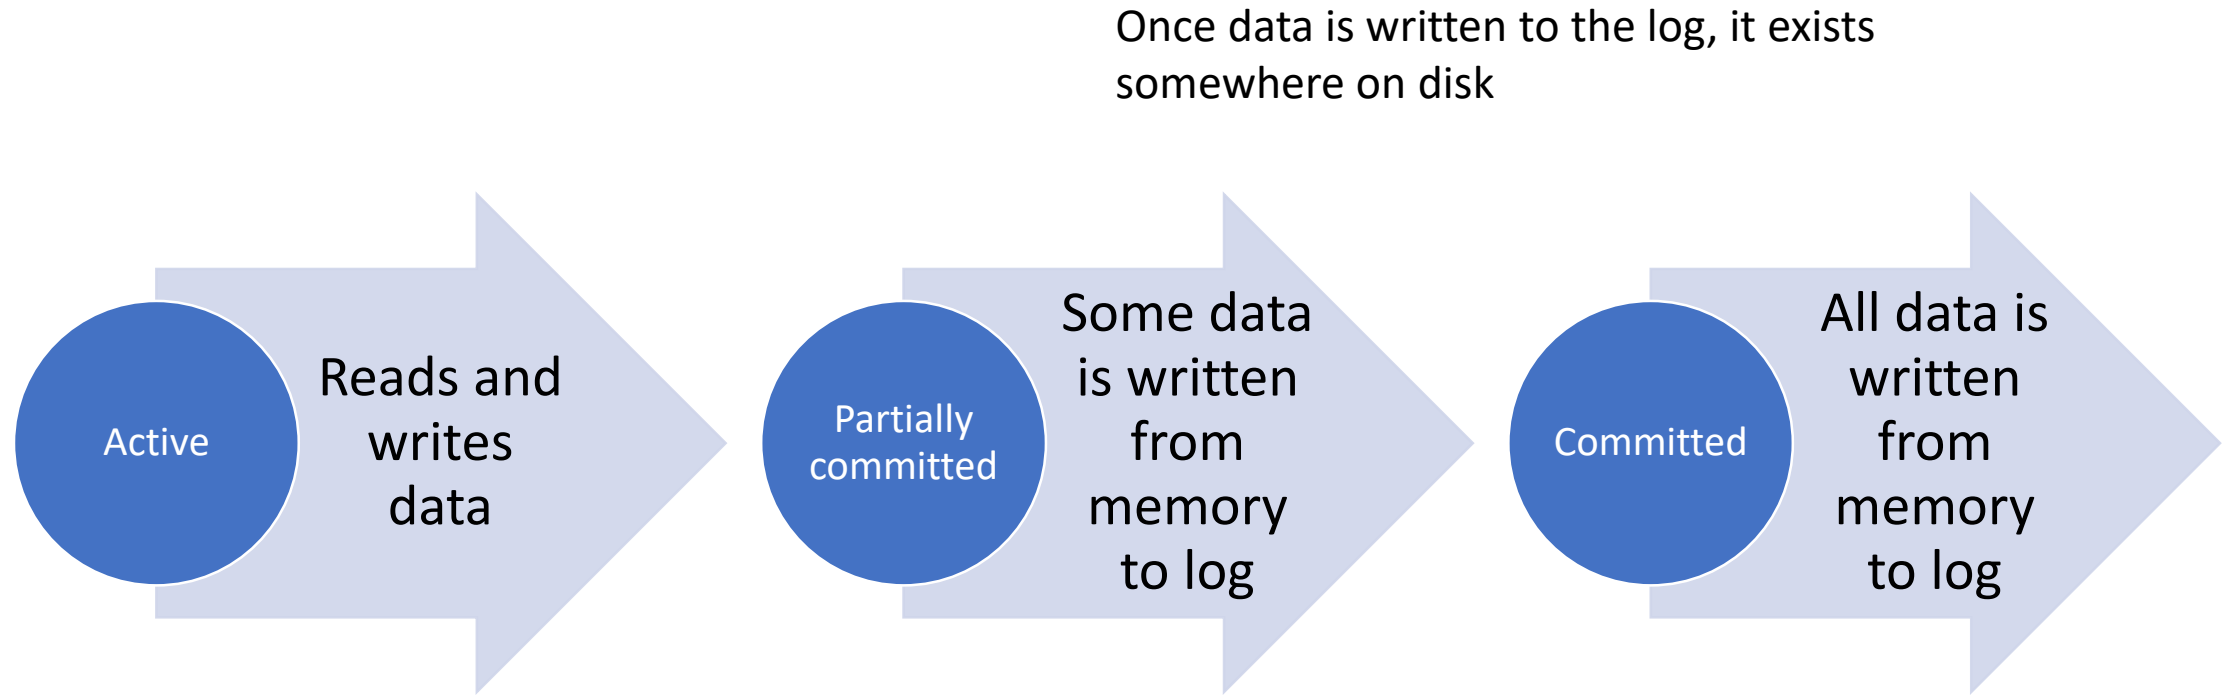


UNIVERSITY OF
CALGARY

Adding more users

- Supporting more users means that databases need to support concurrent transactions
 - Writing data
 - ensure that each item that is to be written to the database is written to disk
 - Reading data
 - ensure that each read operation is accurate
 - Updating data
 - ensure that data is read accurately, then changed in a consistent way

The lifecycle of a transaction



Transactions may be aborted by the application (user) or otherwise fail

Typical problems

- Lost update
 - one interleaved update overwrites another
- Temporary update (dirty reads)
 - a failed transaction does not roll back its data before another operation reads an update that is now invalid
- Incorrect summary
 - aggregation operation reads some database items before they are updated by another operation, and reads some database items after they are updated by that operation
- Unrepeatable reads
 - successive reads by a transaction results in different values, because another operation has updated data between reads

Isolation levels

- Indicates to the DBMS how tolerant of consistency errors you may be
- Generally set at a database-wide level, but some platforms allow per-transaction isolation levels.

SQL standard isolation levels are:

READ UNCOMMITTED	Any consistency errors
READ COMMITTED	No dirty reads
REPEATABLE READ	No dirty or phantom reads
SERIALIZABLE	No consistency errors

Controlling transactions

- To manage which transactions have access to specific parts of the database, objects called **locks** are assigned to individual data items
 - only one lock per object
- Different locking schemes are possible
 - **binary locks**: items are either locked or unlocked
 - **shared/exclusive (read/write) locks**: shared locks allow for reading, exclusive locks must be used for writing
- **Two phase locking** protocols insure that transactions can be serializable
 - **Growing**: transactions can acquire locks
 - **Shrinking**: transactions can release locks

Deadlocks and lock escalation

- The DBMS also monitors when data items are **in contention**
 - Deadlocks: When transaction A has a lock for item X and needs a lock for item Y, but transaction B has a lock for item Y and needs a lock for item X
 - the system will cause one of these transactions to fail to ensure progress
 - items which are in high demand are said to be in contention
- To reduce contention, some platforms will allow for different granularities of lock
 - row level, block(page) level, table-level
 - if many items are in contention, then the granularity level can be changed to reduce demand

In MySQL

- Explore the performance_schema database
 - What is the isolation level? (look at the global_variable table)
 - Is table-level locking allowed? (look at the global_variable table)
 - Can you find anything interesting in the global_status table?)
- Hint: You don't have to switch databases to access tables from another schema. Try:

```
SELECT count(*) FROM performance_schema.global_variables;
```


Adding more data


- Adding more tables
 - How do we use more complex schemata in interesting ways?
 - Can we combine data from different databases in meaningful ways?
- Adding more records
 - What is the impact of big data on databases?
 - Consider relational databases
 - cost of joins
 - cost of constraints
 - cost of transaction management

Adding more computing power and disk


- More complexity in your system architecture results in more complex logical and physical design
 - How do you spread data across different disks and CPUs?
- Consider if you have a single database server, and you're able to add more computing power and disk
 - Naively, you might treat the entire system as a single processing unit
 - CPUs share the same disk and the same memory area
 - A single query might be assigned (in turn) to different CPUs to process

Different approaches to processing

Symmetric multi-processing (SMP)



Massively Parallel Processing (MPP)



Adding more servers

- Consider how to spread data (and tasks) amongst each node
- Heterogeneous vs Homogeneous nodes
 - Are nodes identical, or are they different?
- Physical location
 - are your nodes located close together, or geographically far apart?
 - **Latency** (the time it takes to transmit data between nodes) quickly becomes a factor

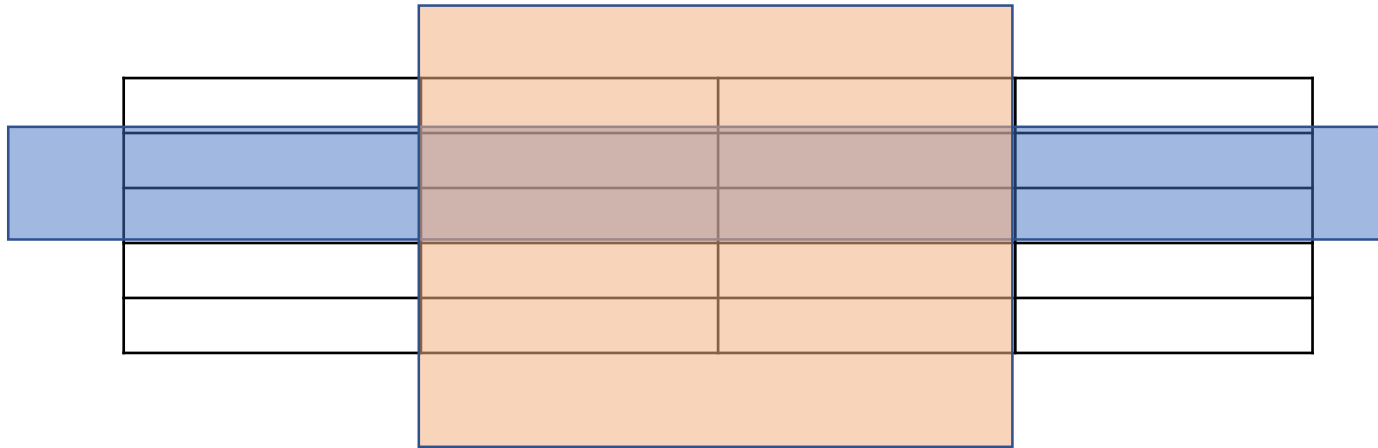
How should clusters share resources?

- There are three typical styles for resource sharing in a distributed database management system

	Memory	Disk
Shared all	Shared	Shared
Shared disk	Not shared	Shared
Shared nothing	Not shared	Not shared

Partitioning

- How should data (and operations) be distributed amongst each node?

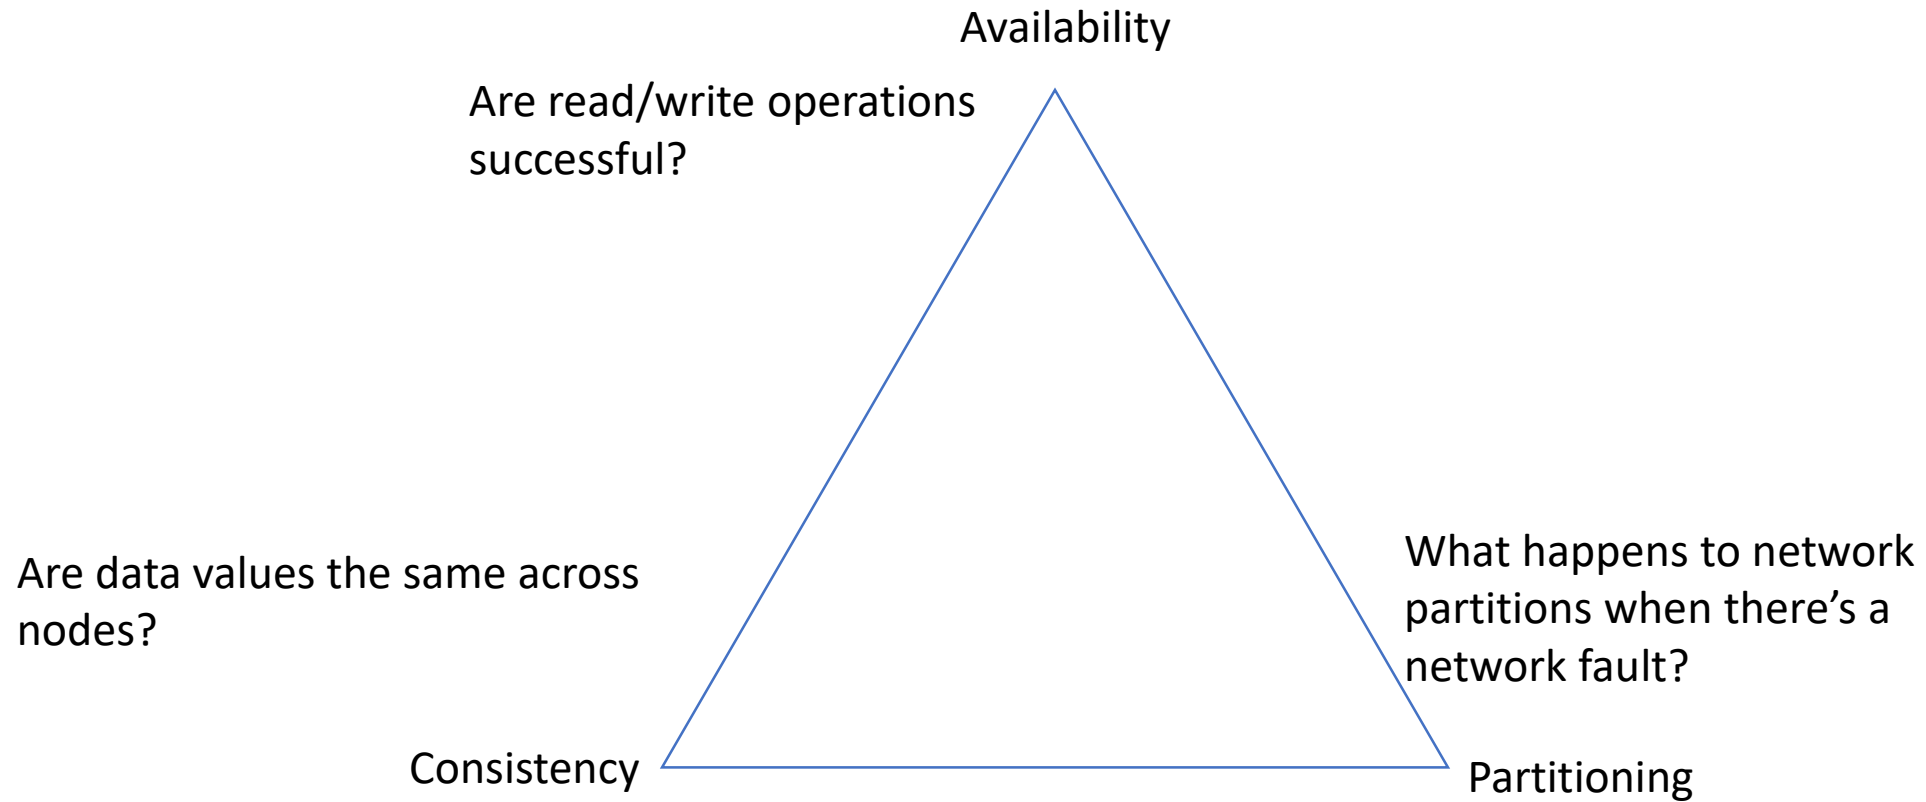


- Horizontally (each node receives a subset of records)
- Vertically (each node receives a piece of each record)
- Often referred to as **sharding** in NoSQL systems

NoSQL and Distribution

- Because NoSQL database management is much simpler, they can be easier to manage at scale across many nodes. Consider the following:
 - Joins
 - Constraints
 - Transactions
- NoSQL databases provide more flexibility when nodes are distributed across a network

The CAP Theorem



The coming weeks

	Topic	What we scale
November 25	Data Mining and Warehousing	Breadth and depth of data
November 27	Big Data	Adding more servers
December 2	Specialty databases	Kinds of data
★ December 4	Cloud computing	Networks

USRIs