

CLASSIFICATION



FOR LATER
DOWNLOAD THE
“Exercise - Classification.zip”
NOTEBOOK & DATASET
FROM THE COURSE SITE

CLASSIFICATION

Classification and Types of Classifiers

Some simple classification algorithms

- k-nearest neighbors (kNN)

- Naïve Bayes

Data Story Presentations

Test/Train & Cross-Validation Procedures

Evaluation

A Few More Clustering Techniques

MACHINE LEARNING APPROACHES

	continuous	categorical
supervised	regression	classification
unsupervised	dimension reduction	clustering





CLUSTERING VS. CLASSIFICATION

Unsupervised

We think there are categories or groupings, but we **don't know exactly what they are.**

For example – Different types of people, countries, animals, etc.

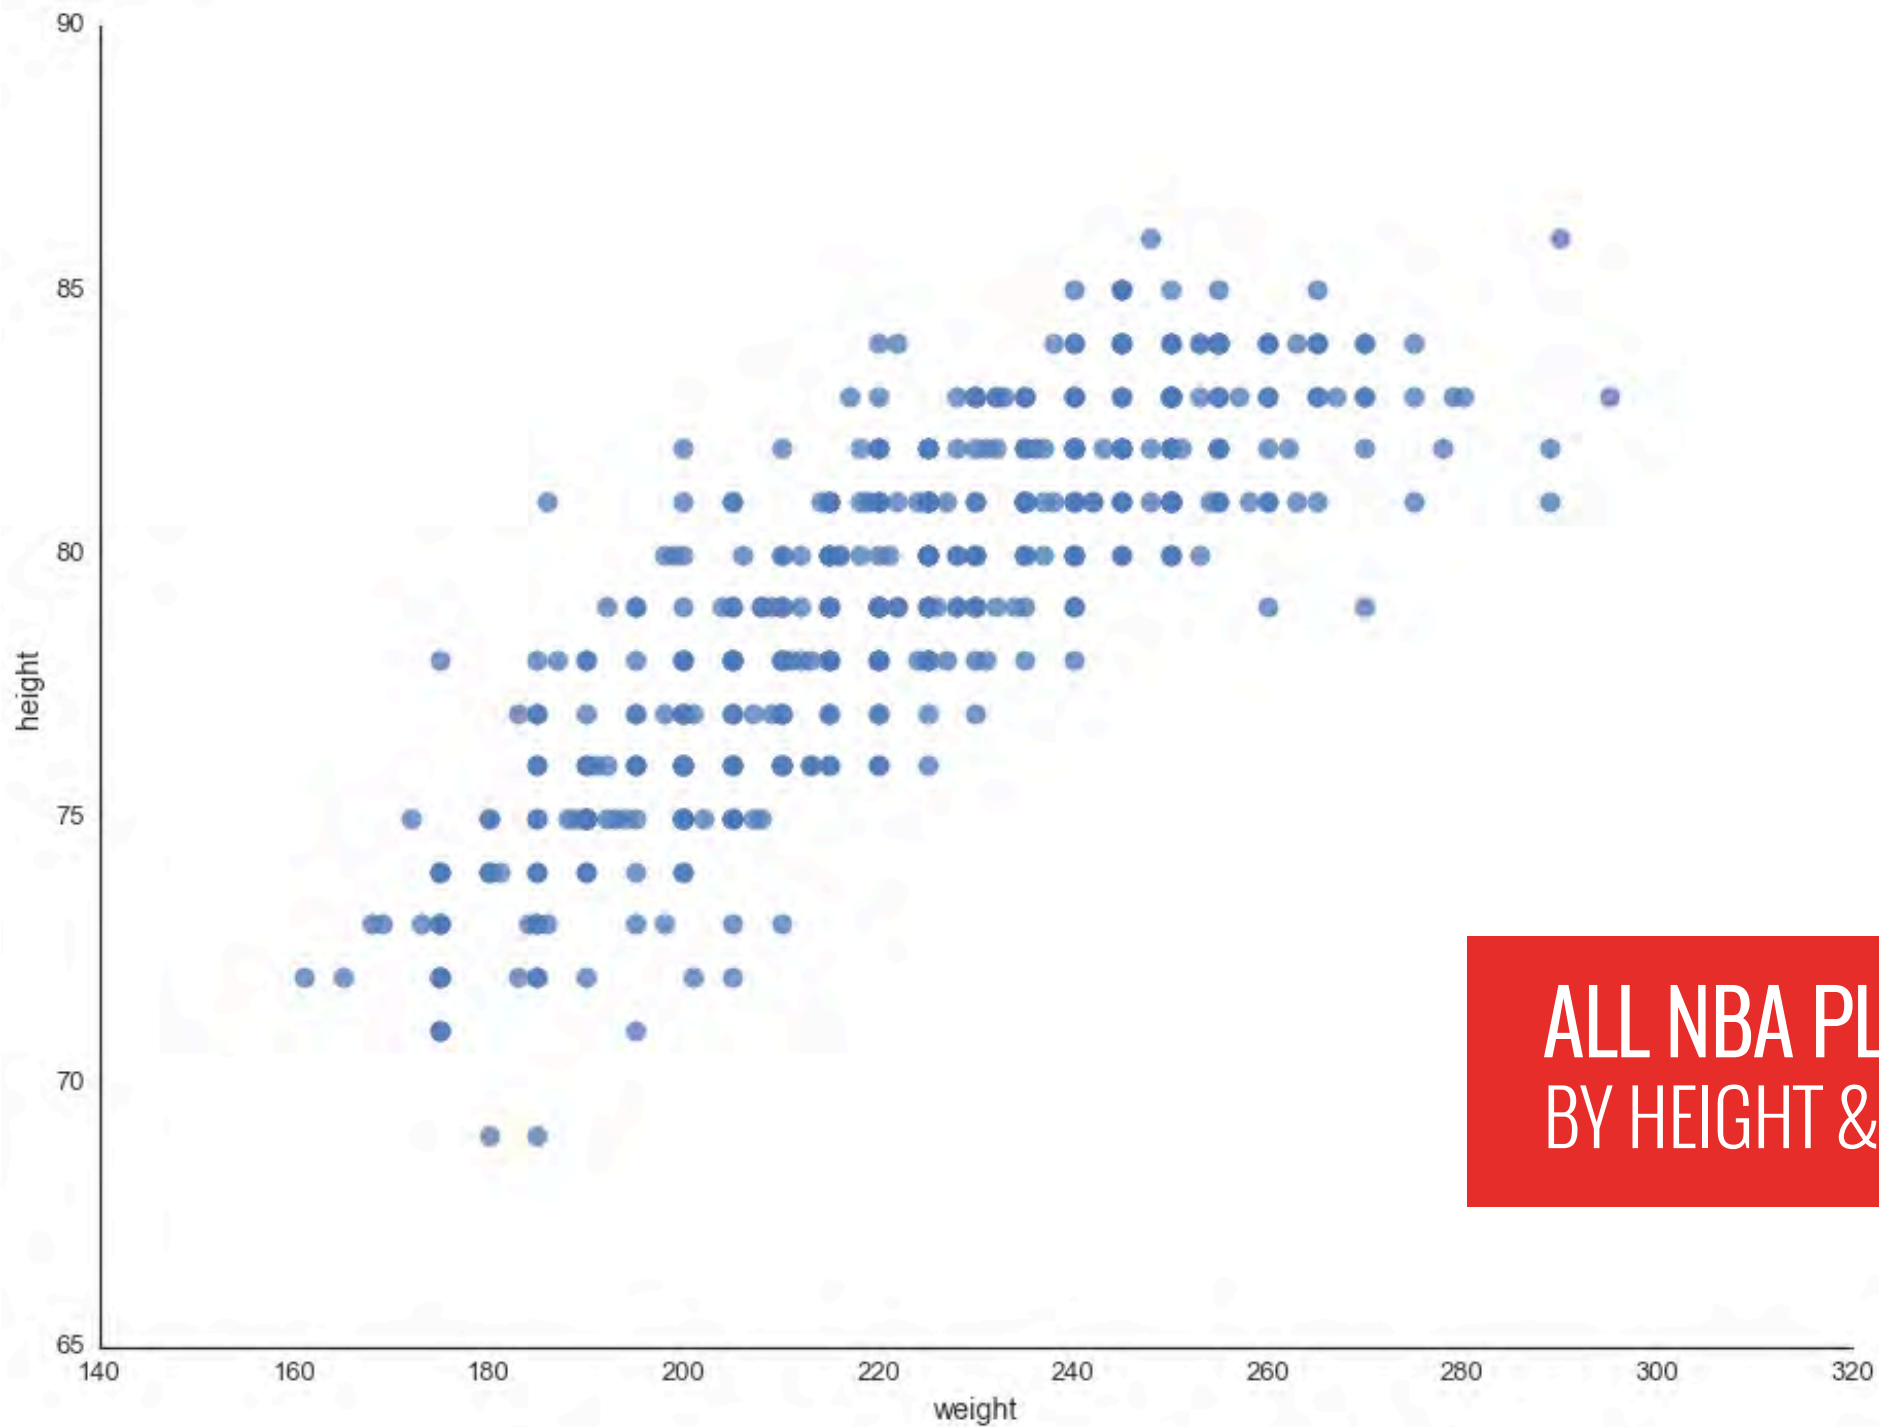
Often supervised (or semi-supervised)

We know specific categories and want to apply them:

Spam versus Not-spam

English versus Spanish versus French

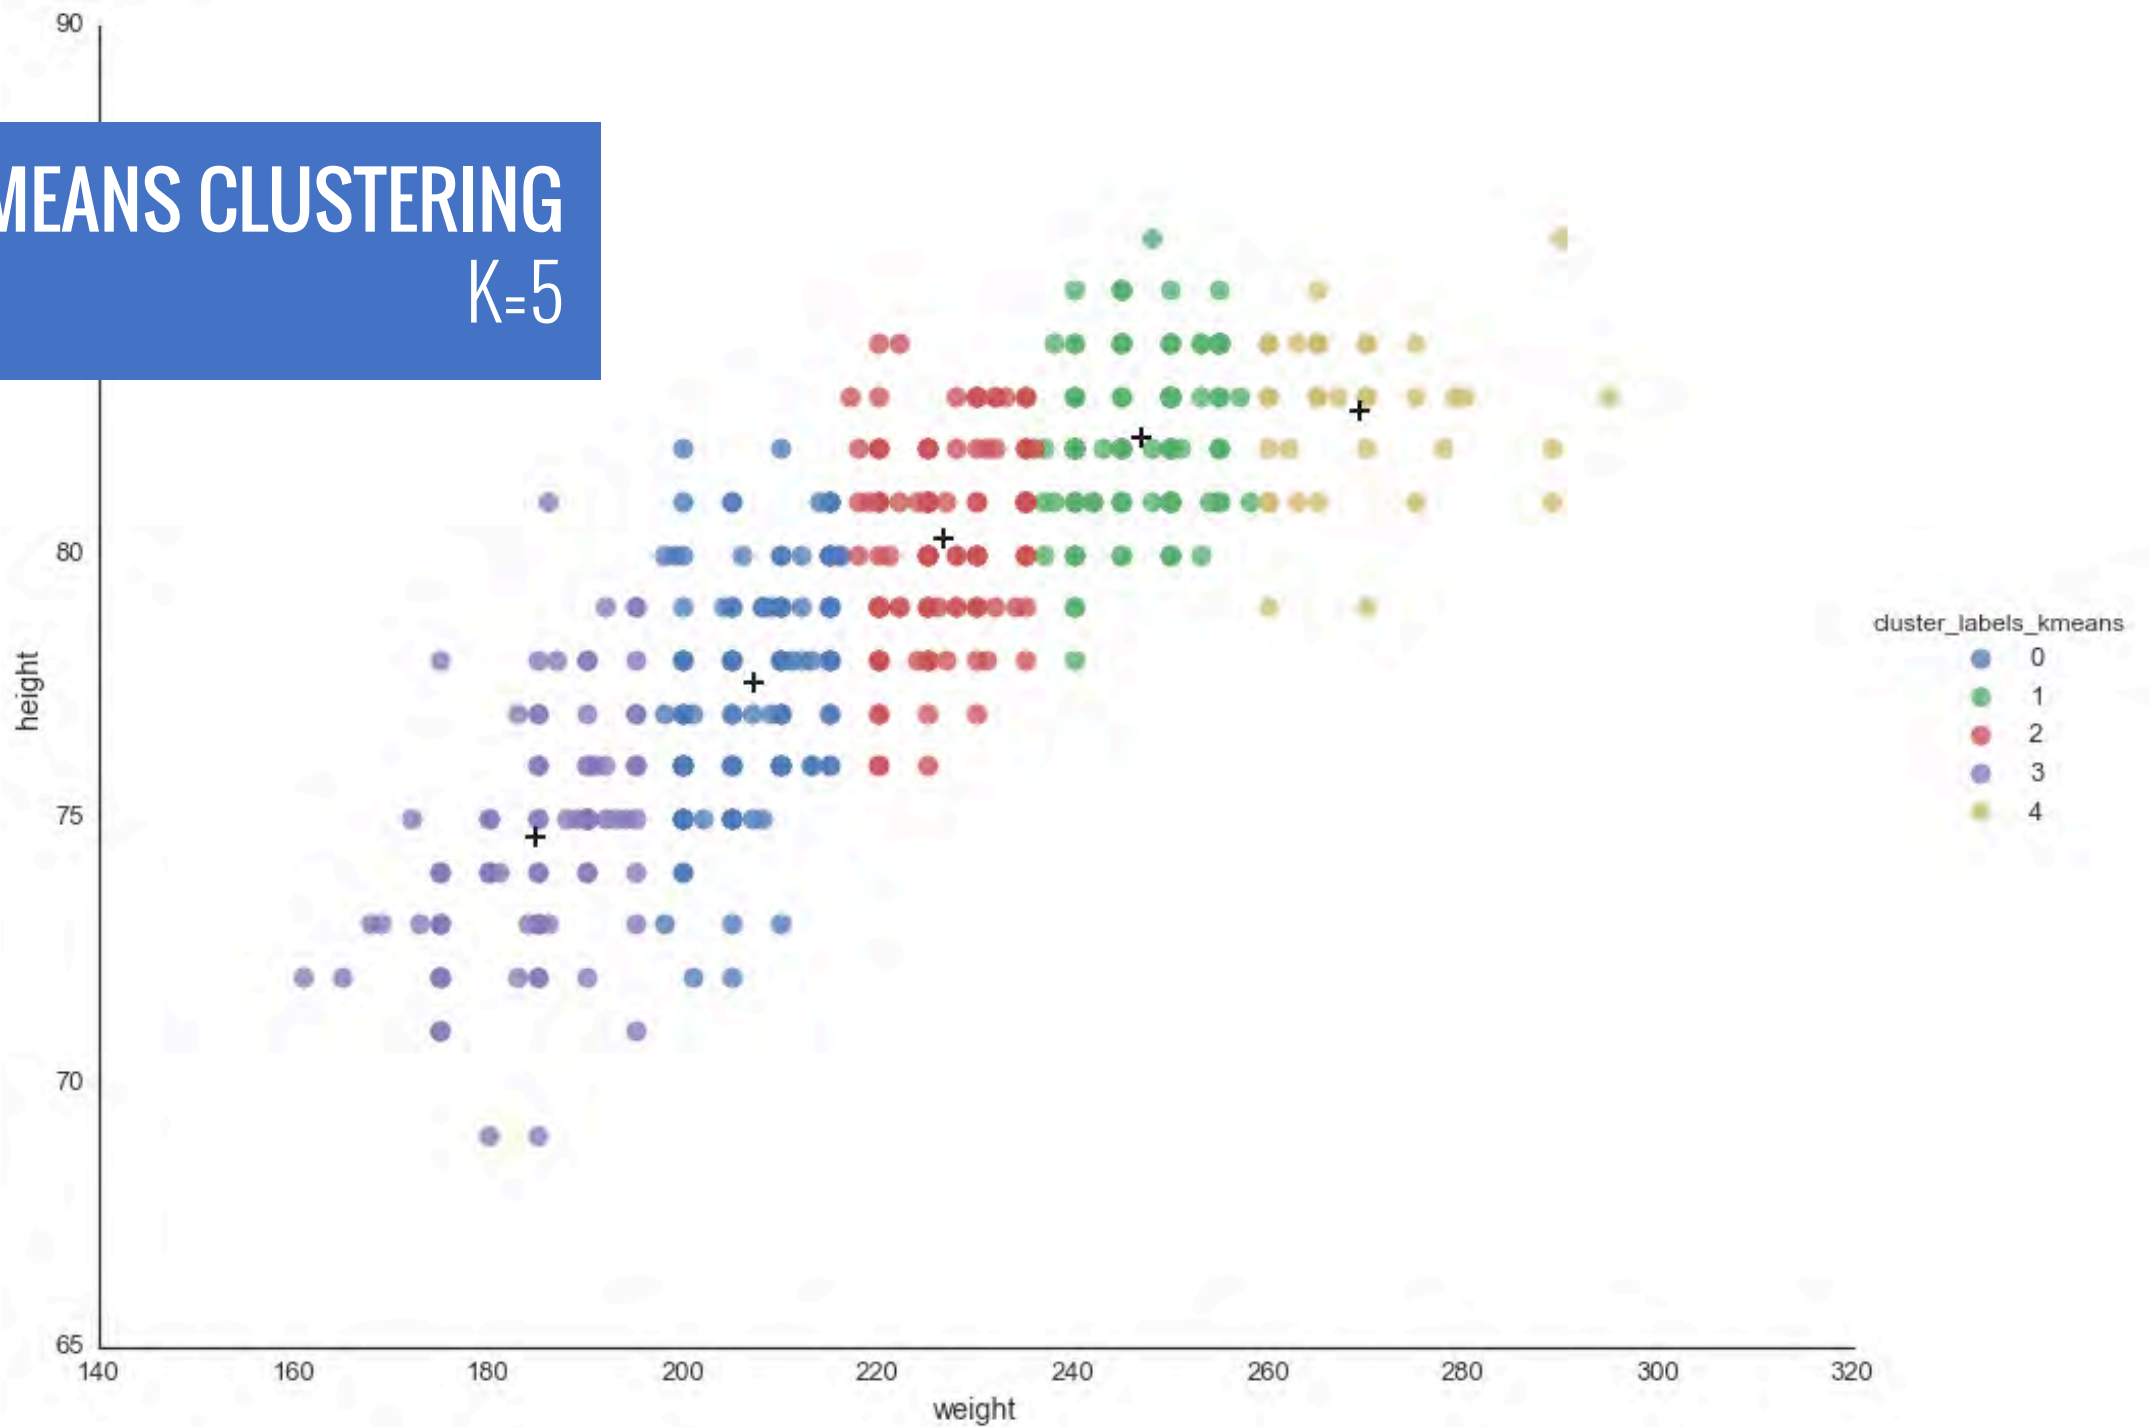
Pop versus Classical

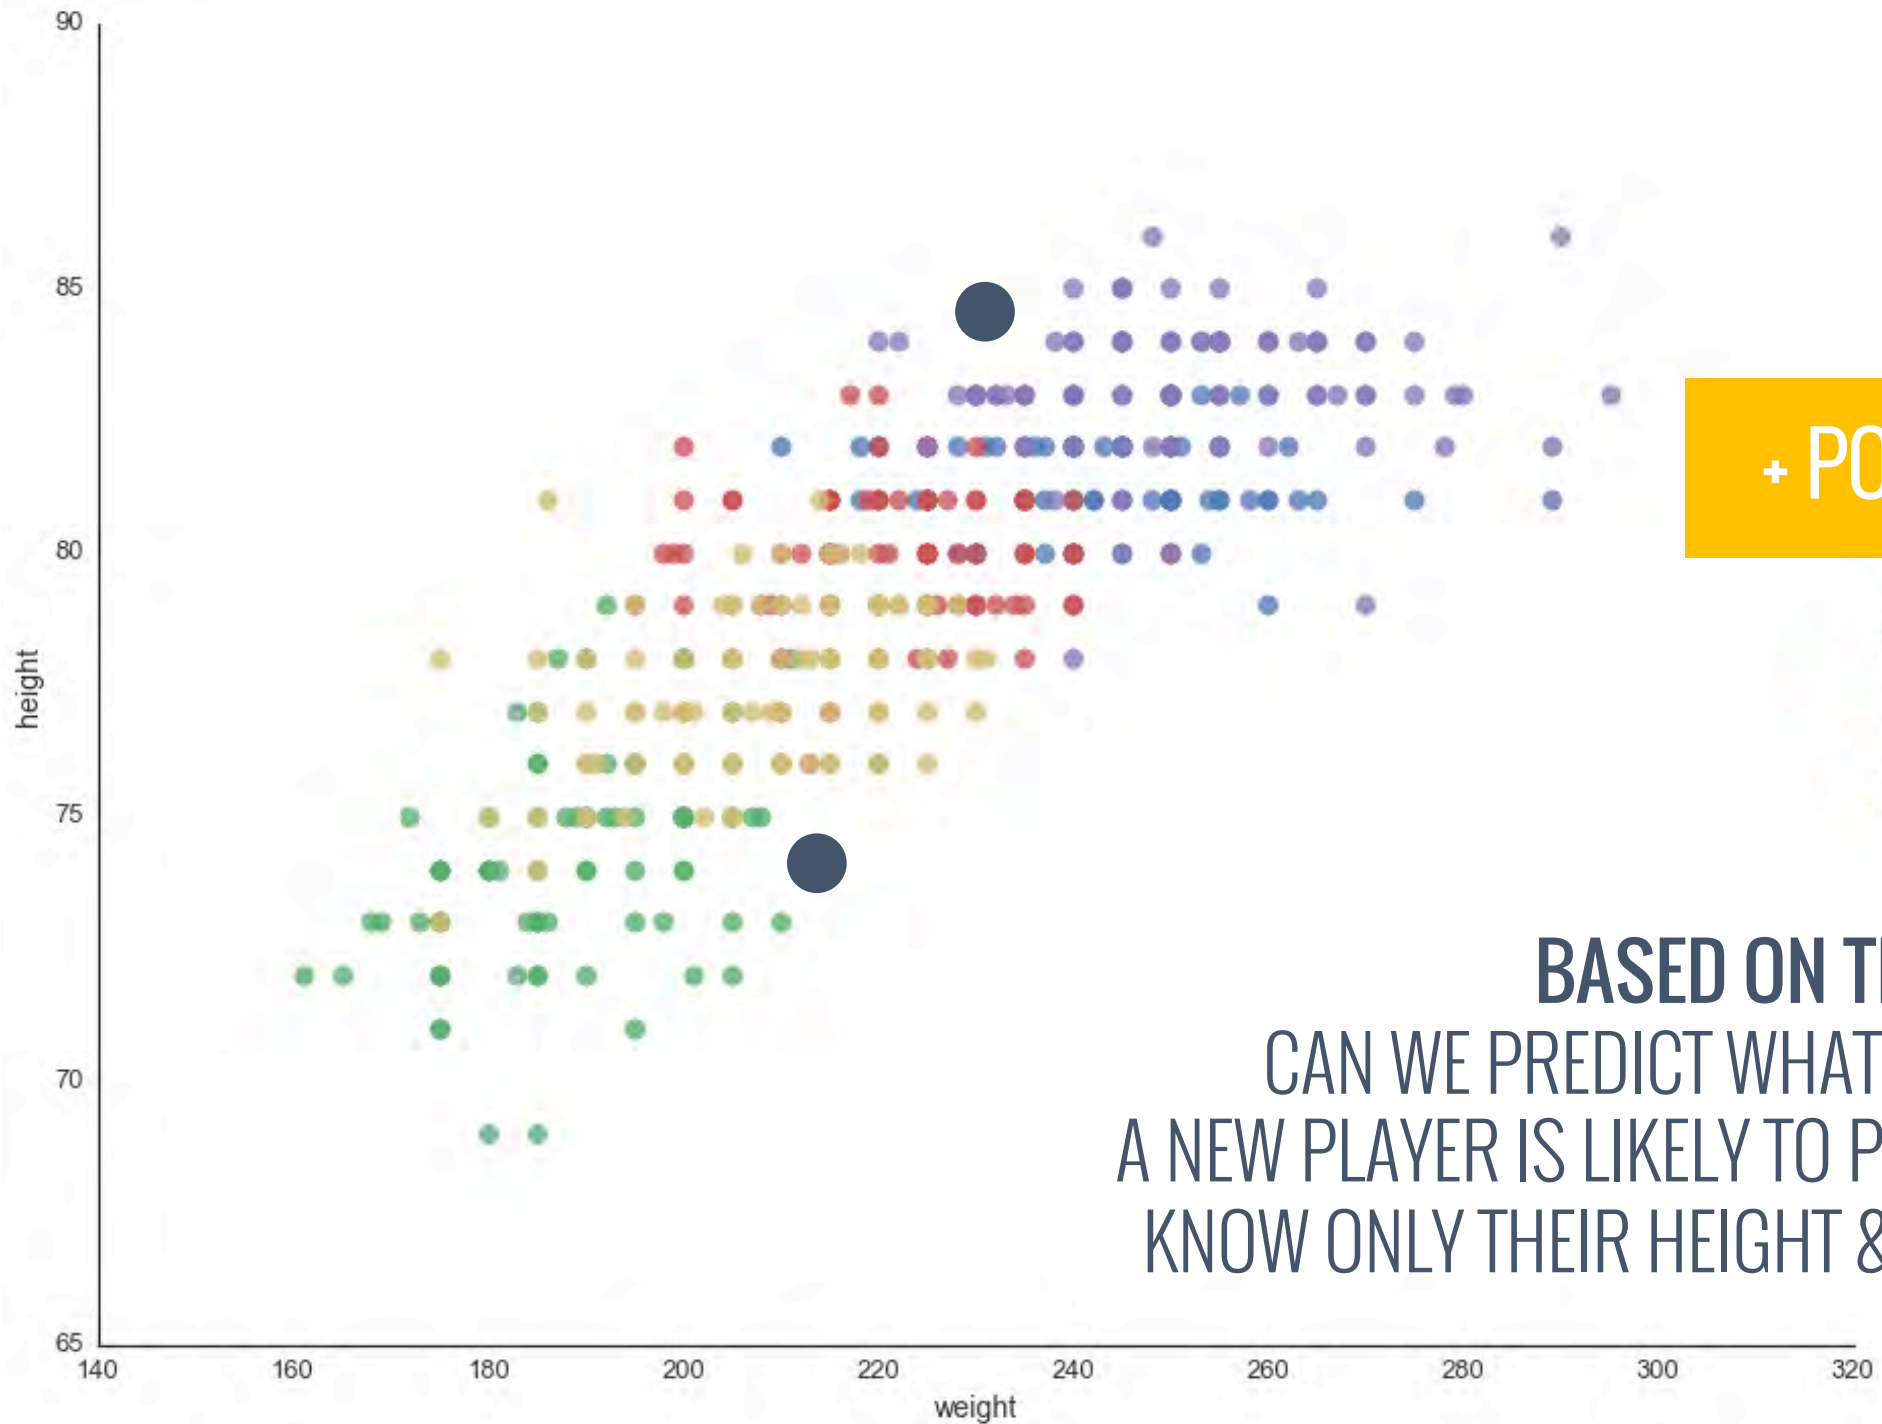


ALL NBA PLAYERS
BY HEIGHT & WEIGHT

K-MEANS CLUSTERING

K=5





+ POSITION

BASED ON THIS DATA
CAN WE PREDICT WHAT POSITION
A NEW PLAYER IS LIKELY TO PLAY, IF WE
KNOW ONLY THEIR HEIGHT & WEIGHT?

CLUSTERING VS. CLASSIFICATION

Clustering

Tries to separate groups by using (dis)similarity


Is player 1 like player 2 but unlike player 3?

Classification

Tries to find important features and weights

What features distinguish a point guard from a center?

CLASSIFICATION PIPELINE



GETTING LABELS

USUALLY THIS IS DATA YOU DON'T HAVE
IF YOU DID, YOU WOULDN'T NEED A CLASSIFIER

The painful part

- Often human labor

 - Hire a bunch of out of work musicians to label your music for you

 - Use Mechanical Turk

Can distribute the pain

- Every time an individual hits the “spam” button in gmail

Can infer labels

- Predict gender by writing style

- Find articles with a byline → guess gender by name (database of names) → tag article by guessed gender

- Noisy, but might be mostly right

Can generate

- Can sometimes generate artificial datasets

 - Make images of cats and dogs in different poses

- Since we generated, we know the labels

TRADEOFFS

How much data do I need?


How hard/expensive is it to label?

How accurate do I need the labels to be?

How much will the “experts” agree?

...

CLASSIFICATION PIPELINE



FEATURES

Same idea as in clustering

Some set of “descriptions” for an object

Explicit: Petal length, player height, miles per gallon, number of times word V1@gra is seen in text, shares sold last period

Inferred/calculated: average rate of change in shares sold in last month

FEATURE ENGINEERING

Could fill an entire course by itself...

Some rules/suggestions:

- If you have lots of features you need **lots of examples**

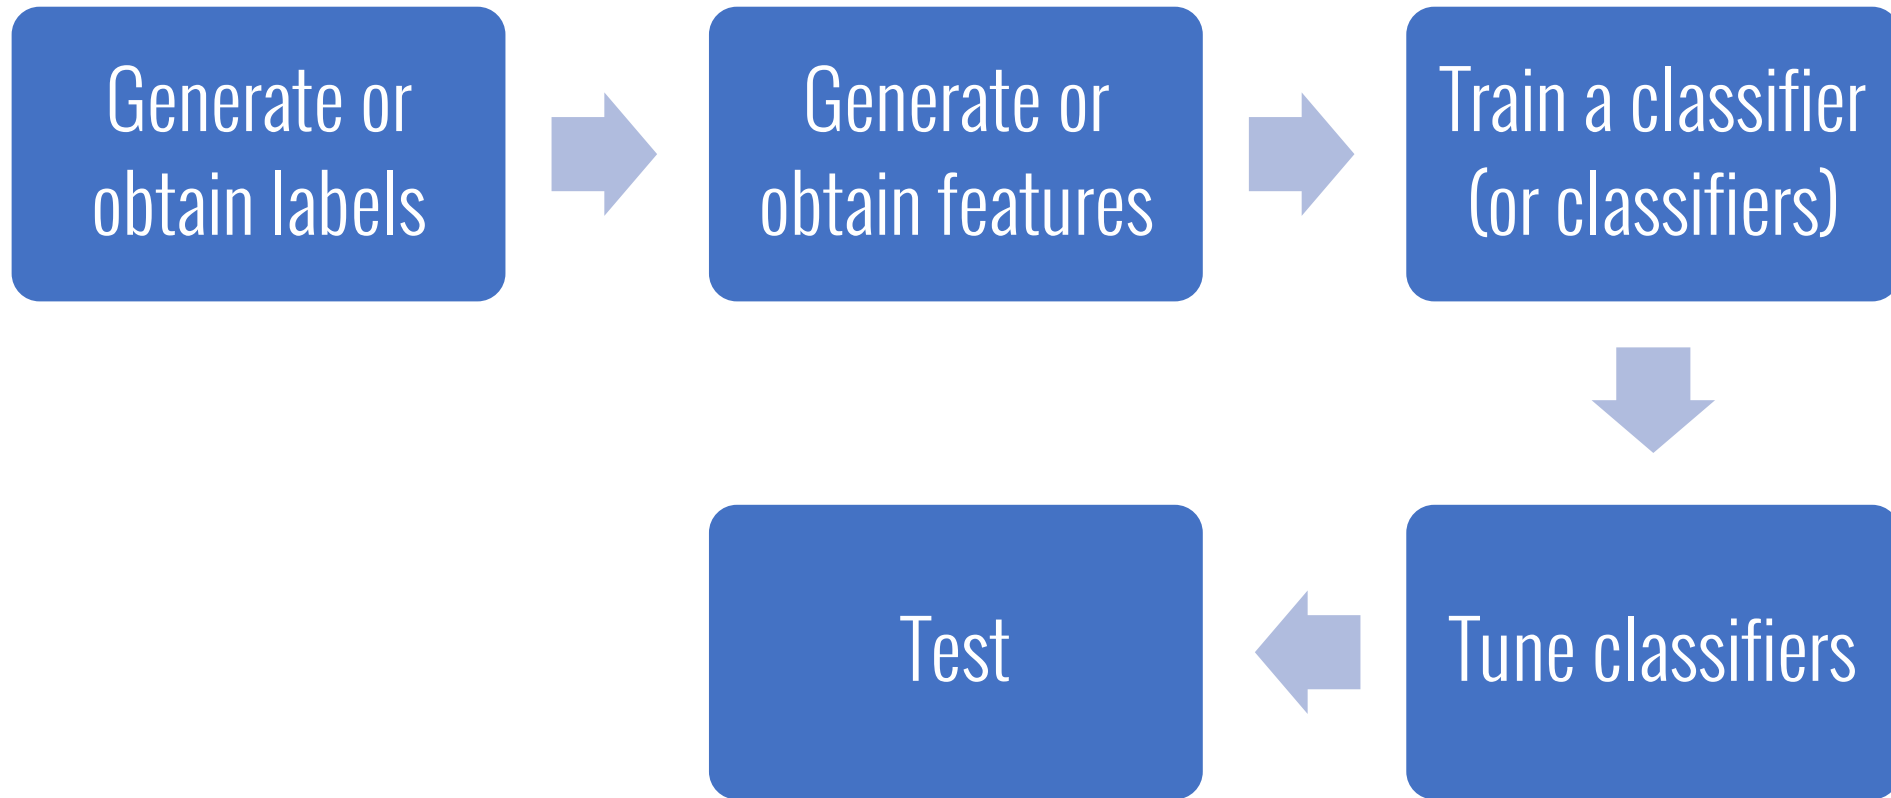
- Use EDA techniques!**

 - Look at **distributions**

 - (is MPG for sports cars obviously different than MPG for station wagons?)

- Start with the **most discriminative features**

CLASSIFICATION PIPELINE



SOME BASIC/POPULAR CLASSIFIERS

kNN (k-Nearest-Neighbor)

Naïve Bayes

Logistic Regression

Decision Trees

Random Forests

SVM (Support-Vector Machines)


Neural Nets (“Deep Learning”)

K-NEAREST NEIGHBOR

(Classification via similarity.)

PICK THE CLOSEST k POINTS

$K=5$ (find closest 5 points)
Pick the majority class



K-NEAREST NEIGHBOR CLASSIFIER

Advantages

- No training needed

- Can be applied to **any distance measure** and feature representation

- Empirically **effective**

Disadvantages

- Finding nearest neighbors has **high time complexity**

- Need to keep the **whole training set**

- Imprecise with small numbers of examples**

 - often true in high-dimensional spaces (neighbors aren't similar enough to be trusted)

NAIVE BAYES CLASSIFICATION

(Classification via probabilistic reasoning.)

BAYES THEOREM

$$\textit{Prob}(A \textit{ given } B) = \frac{\textit{Prob} (A \textit{ and } B)}{\textit{Prob}(B)}$$

GIVES THE PROBABILITY OF AN EVENT OCCURRING GIVEN
THAT ANOTHER EVENT HAS ALREADY OCCURRED

BAYES THEOREM

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

GIVES THE PROBABILITY OF AN EVENT OCCURRING GIVEN
THAT ANOTHER EVENT HAS ALREADY OCCURRED

PREDICTING LABELS

POSTERIOR PROBABILITY
(LIKELIHOOD OF LABEL GIVEN FEATURES)

LIKELIHOOD OF
FEATURES GIVEN
LABEL

LABEL PRIOR PROBABILITY

$$P(L|F) = \frac{P(F|L) \times P(L)}{P(F)}$$

FEATURE PRIOR PROBABILITY

THE PROBABILITY OF AN ITEM GETTING A PARTICULAR LABEL (L) GIVEN THAT IT CONTAINS A PARTICULAR FEATURE (F)

MULTIPLE FEATURES

POSTERIOR PROBABILITY
(LIKELIHOOD OF LABEL GIVEN FEATURES)

LIKELIHOOD OF
MULTIPLE FEATURES
GIVEN LABEL

LABEL PRIOR PROBABILITY

$$P(L|F) = \frac{P(f_1|L) \times P(f_2|L) \times \cdots \times P(f_n|L) \times P(L)}{P(F)}$$

FEATURE PRIOR PROBABILITY

ASSUMES (NAIVELY) THAT ALL FEATURES
ARE COMPLETELY INDEPENDENT
BUT CAN STILL BE REALLY EFFECTIVE!

FOR EXAMPLE

POSTERIOR PROBABILITY
(LIKELIHOOD OF LABEL GIVEN FEATURES)

LIKELIHOOD OF
FEATURES GIVEN
LABEL

LABEL PRIOR PROBABILITY

$$P(spam|words) = \frac{P(viagra, rich, \dots friend|spam) \times P(spam)}{P(viagra, rich, \dots friend)}$$

FEATURE PRIOR PROBABILITY

AN EVEN EASIER EXAMPLE - THE GOLF DATASET

Outlook	Temperature	Humidity	Windy	Play Golf
overcast	hot	high	FALSE	yes
overcast	cool	normal	TRUE	yes
overcast	mild	high	TRUE	yes
overcast	hot	normal	FALSE	yes
rainy	mild	high	FALSE	yes
rainy	cool	normal	FALSE	yes
HOW CAN WE USE THIS TO PREDICT WHETHER OR NOT WE'LL PLAY GOLF ON A NEW DAY? (EVEN IF WE HAVEN'T SEEN A PARTICULAR COMBINATION BEFORE)			TRUE	no
			FALSE	yes
			TRUE	no
			FALSE	no
			TRUE	no
			FALSE	no
			FALSE	yes
			FALSE	yes
sunny	mild	normal	TRUE	yes

COMPUTE PROBABILITY TABLES

Outlook					Temperature					Humidity					Wind				
	Yes	No	P(Yes)	P(No)		Yes	No	P(Yes)	P(No)		Yes	No	P(Yes)	P(No)		Yes	No	P(Yes)	P(No)
Sunny	2	3	2/9	3/5	Hot	2	2	2/9	2/5	High	3	4	3/9	4/5	False	6	2	6/9	2/5
Overcast	4	0	4/9	0/5	Mild	4	2	4/9	2/5	Normal	6	1	6/9	1/5	True	3	3	3/9	3/5
Rainy	3	2	3/9	2/5	Cool	3	1	3/9	1/5										
Total	9	5	100%	100%	Total	9	5	100%	100%	Total	9	5	100%	100%	Total	9	5	100%	100%

Play	P / Total	
Yes	9	9/14
No	5	5/14
Total	14	100%

THEN USE THEM TO COMPUTE PROBABILITIES

For a new day that's rainy, cool, high humidity, and windy:

$$P(\text{Yes}/\text{NewDay}) = \frac{P(\text{RainyOutlook}/\text{Yes}) \times P(\text{CoolTemperature}/\text{Yes}) \times P(\text{HighHumidity}/\text{Yes}) \times P(\text{WithWind}/\text{Yes}) \times P(\text{Yes})}{\cancel{P(\text{NewDay})}}$$

$$P(\text{No}/\text{NewDay}) = \frac{P(\text{RainyOutlook}/\text{No}) \times P(\text{CoolTemperature}/\text{No}) \times P(\text{HighHumidity}/\text{No}) \times P(\text{WithWind}/\text{No}) \times P(\text{No})}{\cancel{P(\text{NewDay})}}$$

$$P(\text{Yes}/\text{NewDay}) + P(\text{No}/\text{NewDay}) = 1$$

COMPUTE

$$P(Yes|NewDay) \propto \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{9}{14} \approx 0.0079$$

$$P(No|NewDay) \propto \frac{2}{5} \cdot \frac{1}{5} \cdot \frac{4}{5} \cdot \frac{3}{5} \cdot \frac{5}{14} \approx 0.0137$$

NORMALIZE

$$P(Yes|NewDay) = \frac{0.0079}{0.0079+0.0137} = 0.36$$

$$P(No|NewDay) = \frac{0.0137}{0.0079+0.0137} = 0.63$$

CHOOSE A LABEL

NO GOLF!

NAÏVE BAYESIAN CLASSIFIER

Advantages:

- Easy to implement

- Very efficient

- Good results for many applications

Disadvantages


- Assumes feature independence, and can break when that assumption is seriously violated (highly correlated data sets).

LET'S TRY SOME CLASSIFYING

TO FOLLOW ALONG
DOWNLOAD THE CLASSIFICATION
NOTEBOOK & DATASET FROM THE
COURSE PAGE

kNN (k-NEAREST NEIGHBORS)

K=5 (find closest 5 points)
Pick the majority class



NAÏVE BAYES CLASSIFICATION

$$P(\text{Yes}|\text{NewDay}) = \frac{P(\text{RainyOutlook}|\text{Yes}) \times P(\text{CoolTemperature}|\text{Yes}) \times P(\text{HighHumidity}|\text{Yes}) \times P(\text{WithWind}|\text{Yes}) \times P(\text{Yes})}{\cancel{P(\text{NewDay})}}$$

$$P(\text{No}|\text{NewDay}) = \frac{P(\text{RainyOutlook}|\text{No}) \times P(\text{CoolTemperature}|\text{No}) \times P(\text{HighHumidity}|\text{No}) \times P(\text{WithWind}|\text{No}) \times P(\text{No})}{\cancel{P(\text{NewDay})}}$$

		Outlook					
	Yes	No	P(Yes)	P(No)		Yes	
Sunny	2	3	2/9	3/5	Hot	2	
Overcast	4	0	4/9	0/5	Mild	4	
Rainy	3	2	3/9	2/5	Cool	3	
Total	9	5	100%	100%	Total	9	

COMPUTE

$$P(\text{Yes}|\text{NewDay}) \propto \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} \cdot \frac{9}{14} \approx 0.0079$$

$$P(\text{No}|\text{NewDay}) \propto \frac{2}{5} \cdot \frac{1}{5} \cdot \frac{4}{5} \cdot \frac{3}{5} \cdot \frac{5}{14} \approx 0.0137$$

NORMALIZE

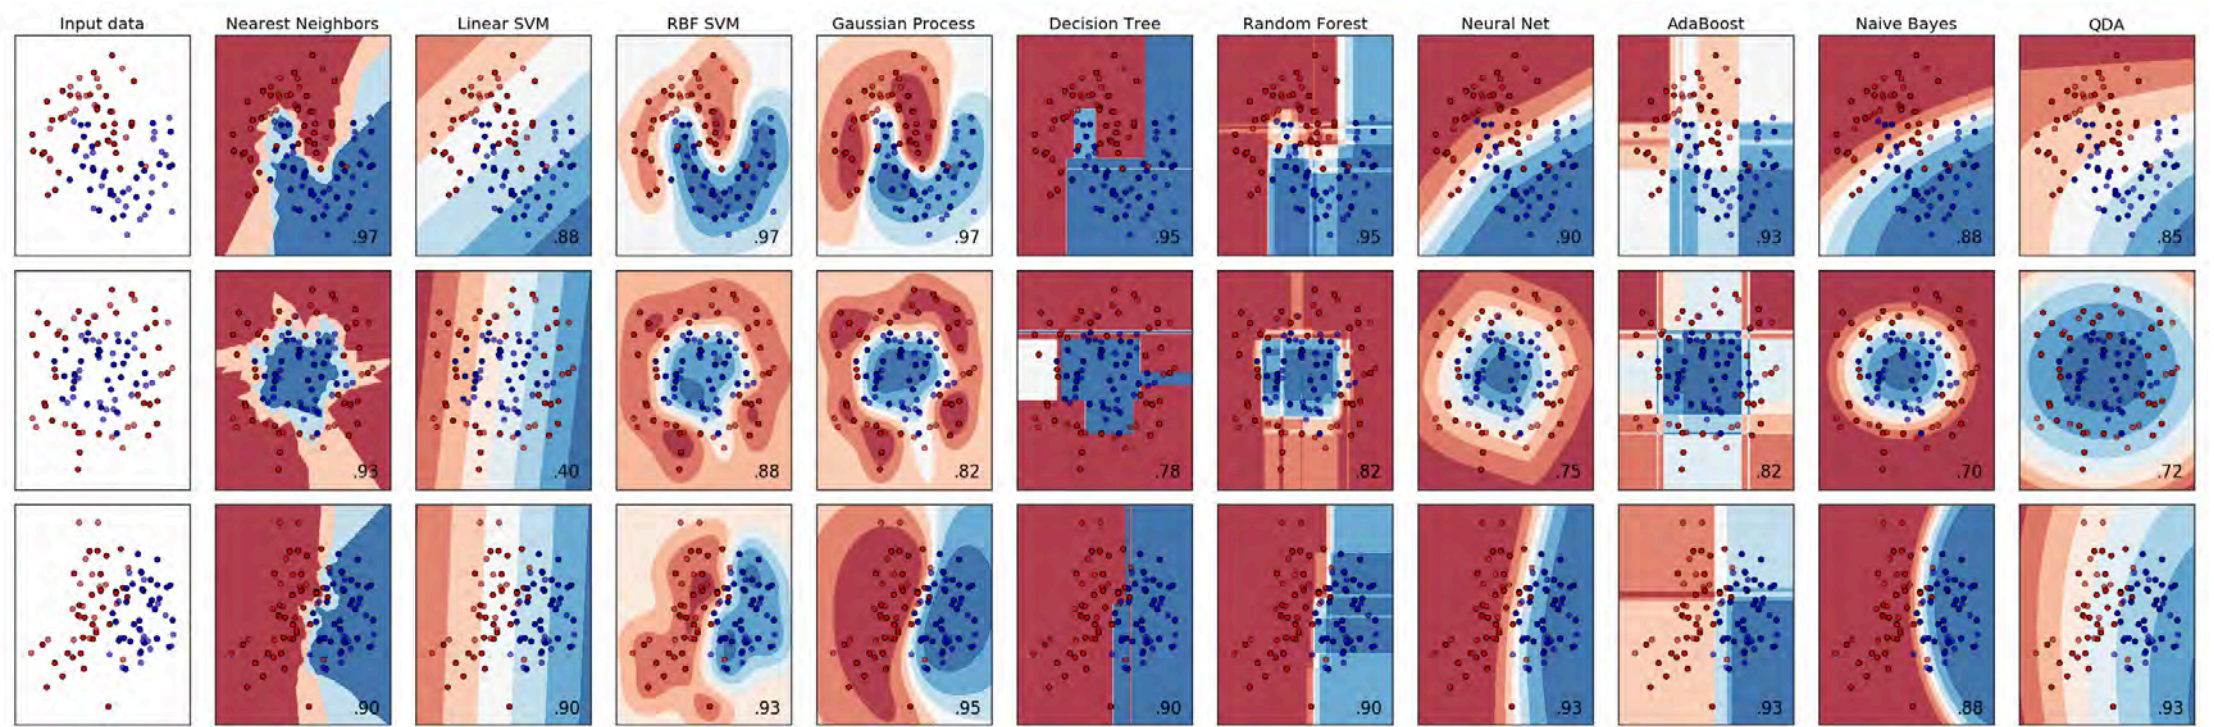
$$P(\text{Yes}|\text{NewDay}) = \frac{0.0079}{0.0079+0.0137} = 0.36$$

$$P(\text{No}|\text{NewDay}) = \frac{0.0137}{0.0079+0.0137} = 0.63$$

THERE ARE MANY OTHER OPTIONS AT YOUR DISPOSAL

Just a few of the classifiers available in scikit-learn:

```
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.gaussian_process.kernels import RBF
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
```



Different classifiers can give dramatically different results.

BUT FIRST...

TRAINING

CLASSIFICATION TRAINING

All our algorithms need to be “trained”

Our algorithm also needs to be “tested”

What happens if we use all our data for training?

“Overfitting!”

CLASSIFICATION TRAINING

First impulse (usually a reasonable one)

Split the data into **train** and **test** (aka “holdout”)

e.g., 70% to train, 30% retained to test

Still need to be careful

You might need to train again!

Some algorithms work best with “**balanced**” training

Too many positive or negative examples can cause problems

Sometimes we don’t have enough data to do this (expensive labels)

Can still overfit

THREE-WAY SPLIT

Training

Validation

Use to tune parameters

Test

After running this one, you're done


CROSS-VALIDATION

An even better option - Try many holdouts

Randomly select test/train data

Pick a different 70% each time


Average scores at the end



EXTENSION: K-FOLD CROSS-VALIDATION

Move a sliding window

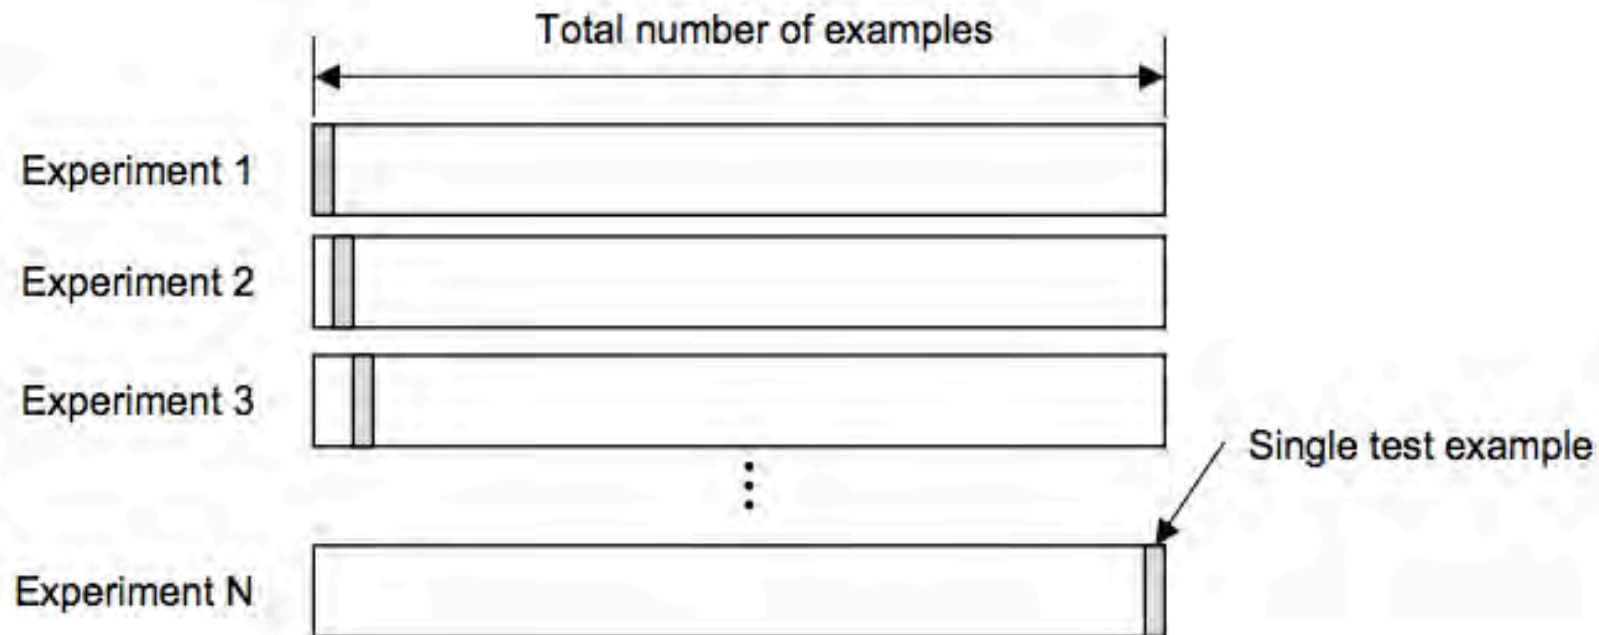
Makes sure that all data is used



EXTENSION: LEAVE-ONE-OUT

Extreme version of k-fold

Remove one at a time, train on all the others



GENERAL GUIDANCE ON FOLDS

Many folds

Accurate but expensive (variance might be high)

More data = get away with less folds

Sparse = use more folds

Common choice = 10

OK...
SO DOES IT WORK?

EVALUATION

Accuracy

Precision

Recall

F1 score

Confusion Matrix

Precision-Recall curve

ROC curve

ACCURACY

Simplest of evaluations

Count the percentage of times we classified correctly.

Algorithm:

- For each item:

 - Compare the label we generated to the ground truth

 - Add 1 if correct, otherwise add 0

- Divide by number of items

PROBLEMS WITH ACCURACY

Think of a spam filter, why might accuracy be the wrong measure?

Hint: What's the distribution of messages? How many spam and how many good?

To win at spam detection (or any unequal task)

– **just guess the majority type.**

If 90% of messages are good: guess that **all are good**

Accuracy is **guaranteed to be 90%**

PRECISION AND RECALL

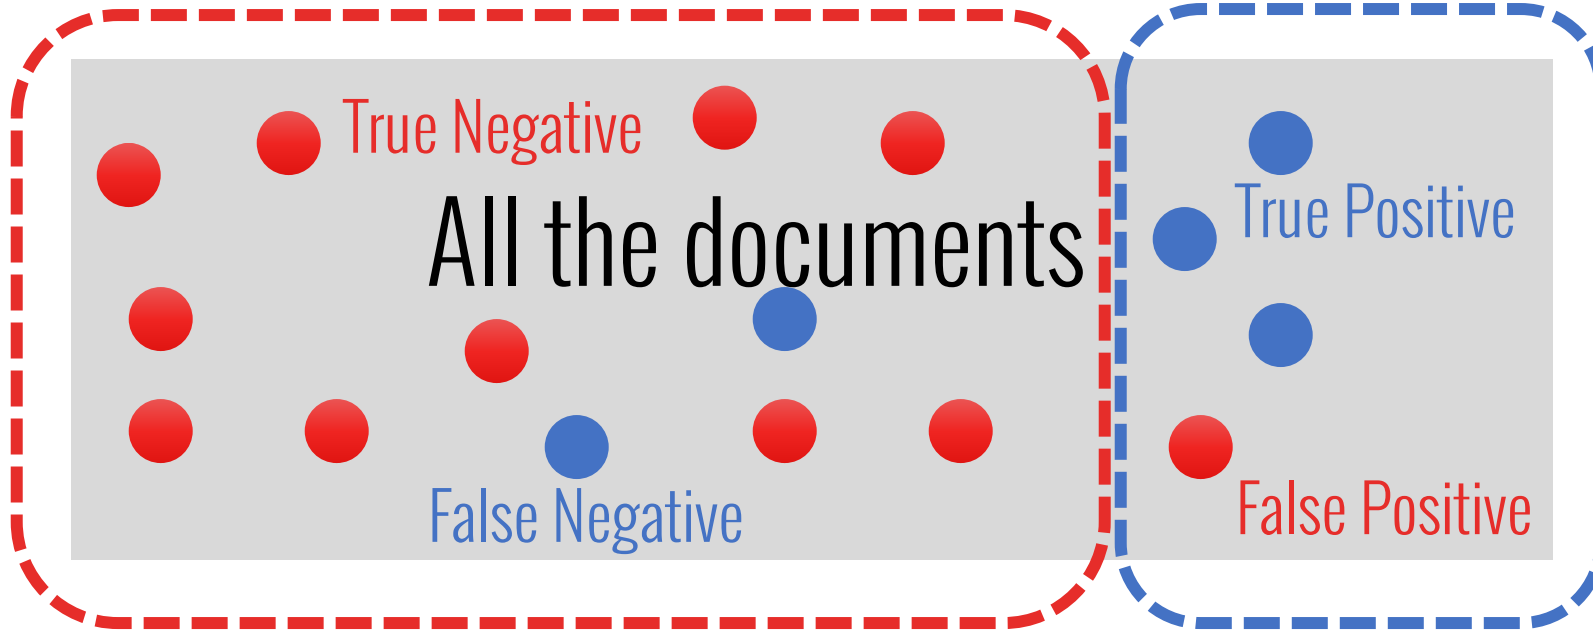
Heavily used in Information Retrieval

Assumption: we have a query “apple”

Our corpus contains documents about apples (positive) and documents about other topics (negative)

● Documents actually about apples

Documents the
system says are
about apples

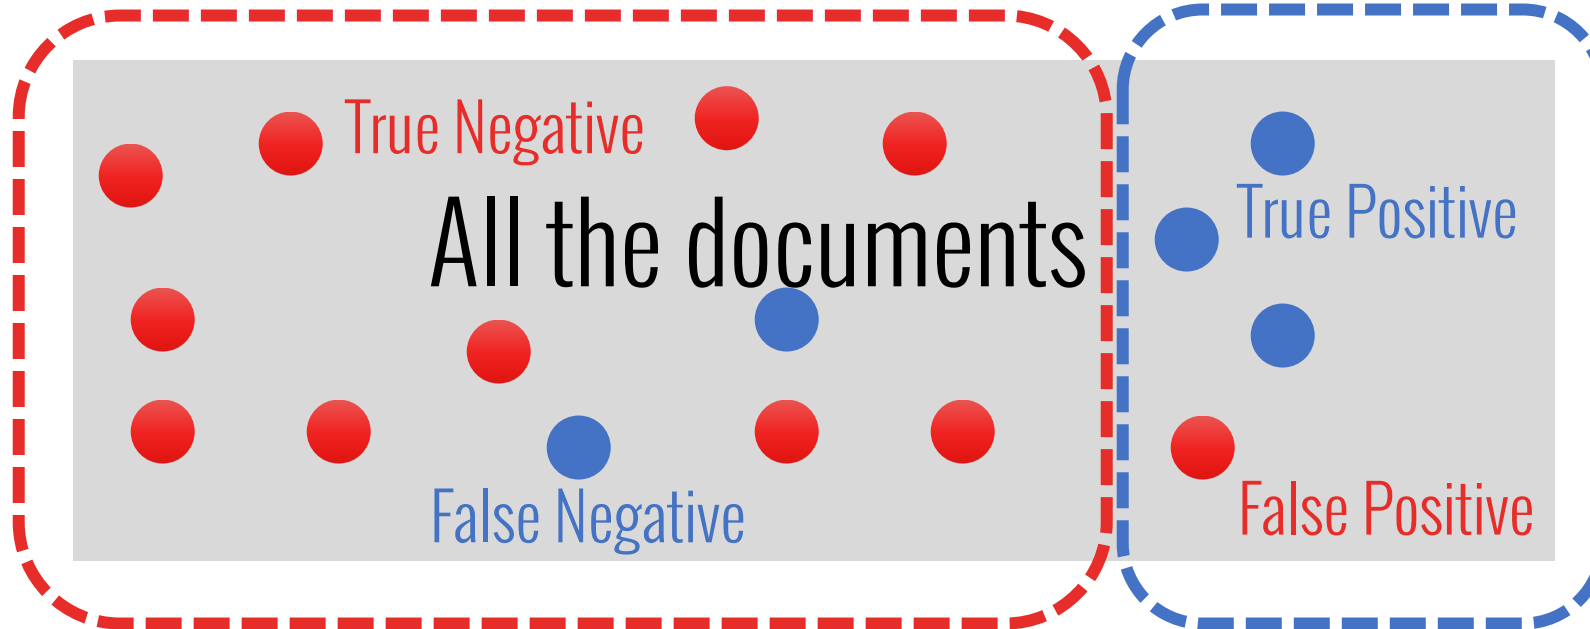


Documents the system
says are **not** about apples

CONFUSION MATRIX (SIMPLE)

	Actually Positive	Actually Negative
Predicted Positive	True Positive	False Positive
Predicted Negative	False Negative	True Negative

● Documents actually about apples



Documents the system
says are **not about apples**

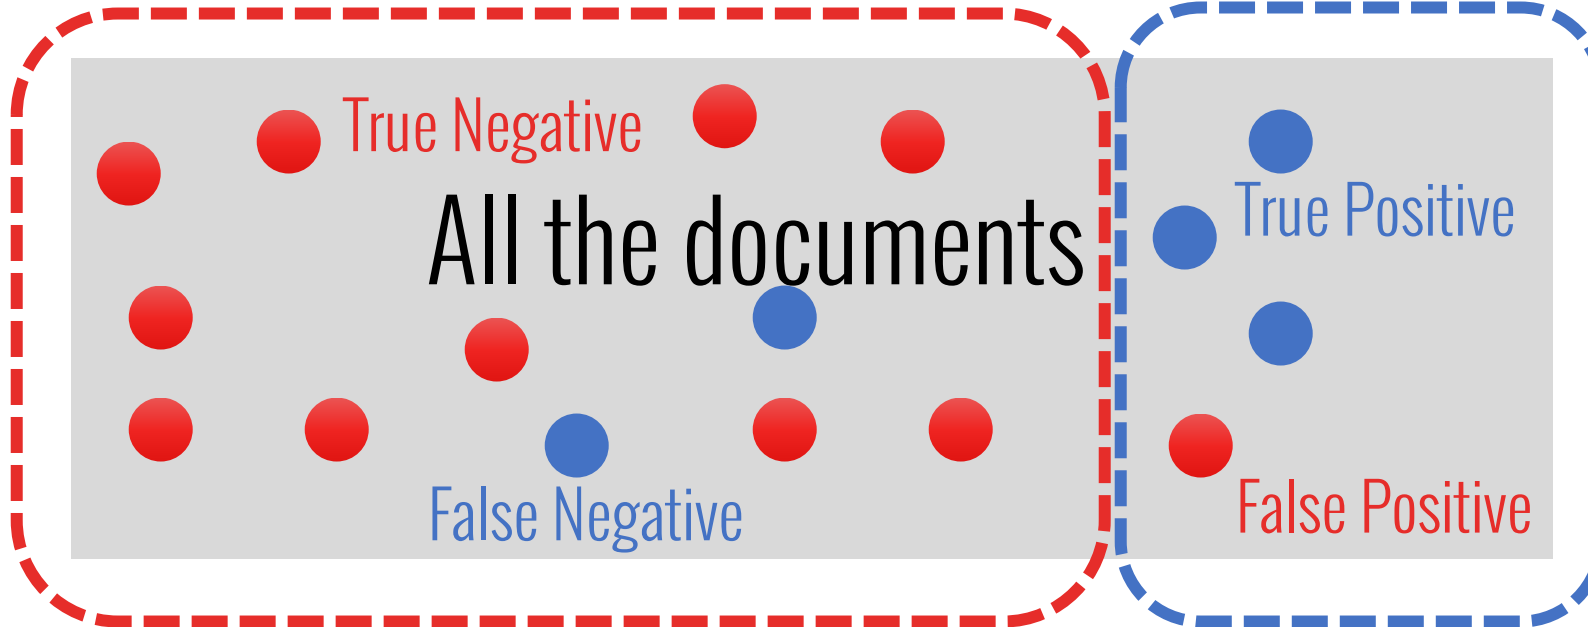
Documents the
system says are
about apples

Precision → % of the items
returned that are true positive

$$TP / (TP + FP)$$

● Documents actually about apples

Documents the
system says are
about apples



Documents the system
says are **not about apples**

Recall → % of the positive
items returned

$$TP / (TP + FN)$$

HOW MANY OF THE POSITIVE
LABELS WERE CORRECT?

Precision: $TP / \text{Predicted positive}$

Recall: $TP / \text{Real positive}$

HOW MANY OF THE REAL
POSITIVES DID WE LABEL?

F1 SCORE

Annoying to deal with two measures

Why not just combine them into one?

F1 Measure: weighted average of precision and recall:

$$F1 = 2 * (Recall * Precision) / (Recall + Precision)$$

Good if you think false positives and false negatives are relatively the same “badness”

FOR MULTI-CLASSIFIERS

If instead of two classes (positive / negative) we have multiple classes (“spam” / “high priority” / “normal”), we can compute precision and recall in several ways.

One option – take class one at a time:


“spam” = positive, all other classes = negative

“high priority” = positive, all other classes = negative

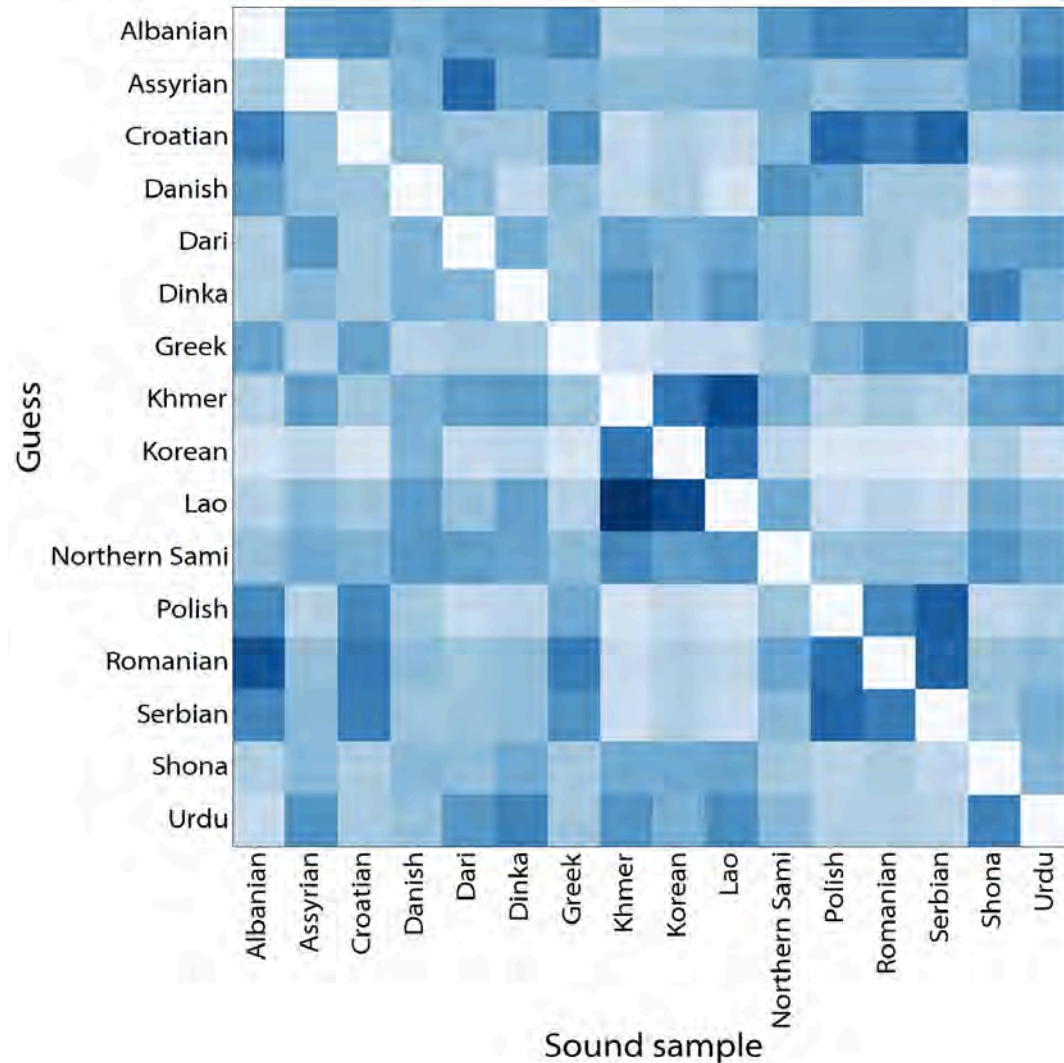
“normal” = positive, all other classes = negative

Calculate *mean* Precision/Recall

CONFUSION MATRIX (MULTI-CLASS)



CONFUSION MATRIX (MULTI-CLASS)



<http://quietlyamused.org/blog/2014/03/12/language-confusion/>


**BACK TO THE
NOTEBOOK!**

DECISION TREES

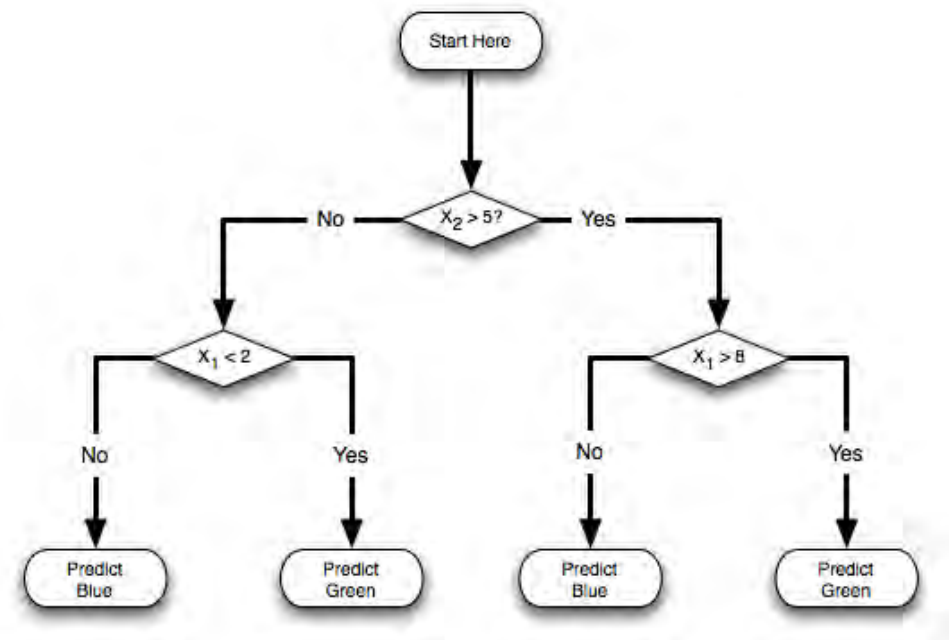
Classification as sets of (usually) binary decisions based on data attributes.


SHOULD I DO LAUNDRY?






REALLY USEFUL IF YOU
CAN CARVE UP THE
ATTRIBUTE SPACE USING
A SET OF IF-THEN RULES







depth = 2



depth = 3



depth = 4



DECISION TREES

Advantages:

Easy to interpret (not all classifiers can be explained)

Prediction process obvious

Can handle mixed data types

INTUITION FOR CONSTRUCTING DTs

Ask the question with the most valuable answer

If I knew the answer to this, how much closer to the solution would I be?

Solutions that divide the space 50/50 are better than solutions that divide the space 2/98.

MEASURING “INFORMATION GAIN”


Use **entropy** or **purity** metrics to assess information gain.

For a set of items S with J classes

Entropy:

$$H(S) = - \sum_{i=1}^J p_i \log_2 p_i$$

p_i = fraction of items in the set with that class



Entropy =
0 for a homogeneous sample
1 for an equally divided one

TO ASSESS THE INFORMATION GAIN OF A SPLIT

$$\text{Gain}(T, a) = \overset{\text{Entropy of Tree}}{H(T)} - \overset{\text{Sum of Entropy of Children after Splitting on Attribute } a}{H(T, a)}$$

Gain will be **high** if a split produces **pure subtrees**.

Can **compute information gain** for all possible splits and then choose the one with the **greatest gain**.



FOR EXAMPLE... HOUSING LOAN DATA

Approved or not

ID	Age	Has_Job	Own_House	Credit_Rating	Class
1	young	false	false	fair	No
2	young	false	false	good	No
3	young	true	false	good	Yes
4	young	true	true	fair	Yes
5	young	false	false	fair	No
6	middle	false	false	fair	No
7	middle	false	false	good	No
8	middle	true	true	good	Yes
9	middle	false	true	excellent	Yes
10	middle	false	true	excellent	Yes
11	old	false	true	excellent	Yes
12	old	false	true	good	Yes
13	old	true	false	good	Yes
14	old	true	false	excellent	Yes
15	old	false	false	fair	No

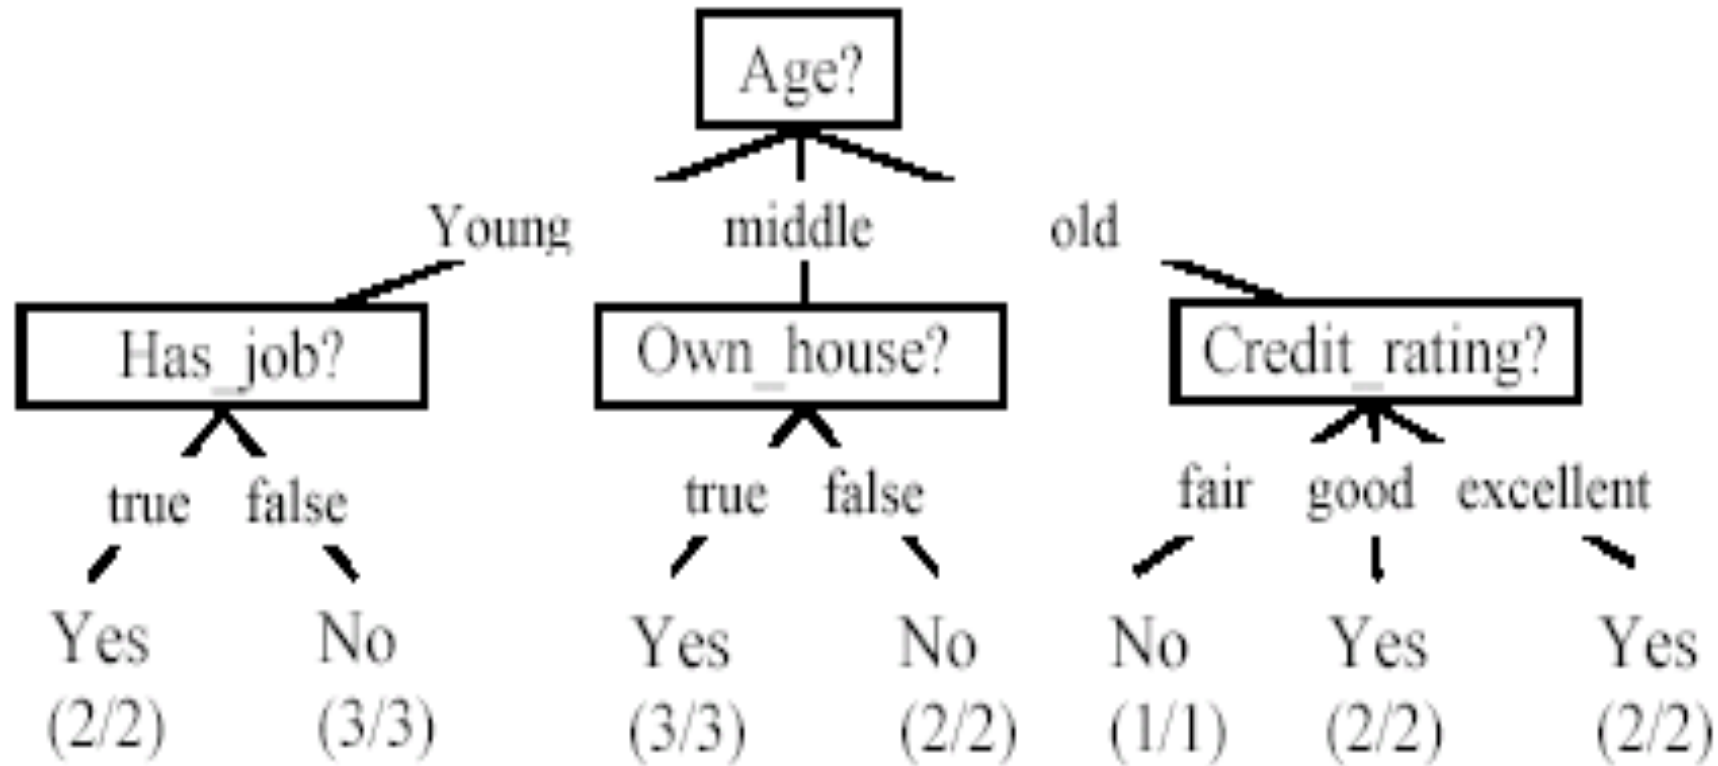
A RECURSIVE ALGORITHM

1. **Split** on the best feature (lowest partition entropy)
2. **Add** a decision node
3. **Repeat** (recursively) with each group of children
4. **Stop** if we hit an entropy threshold or partitions get too small

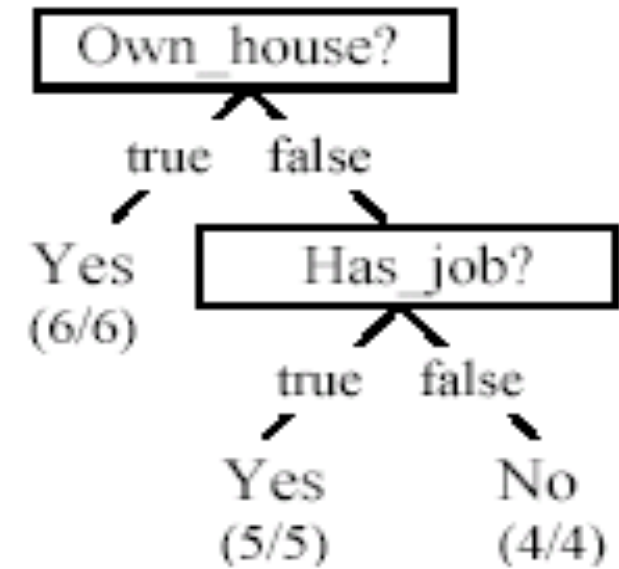
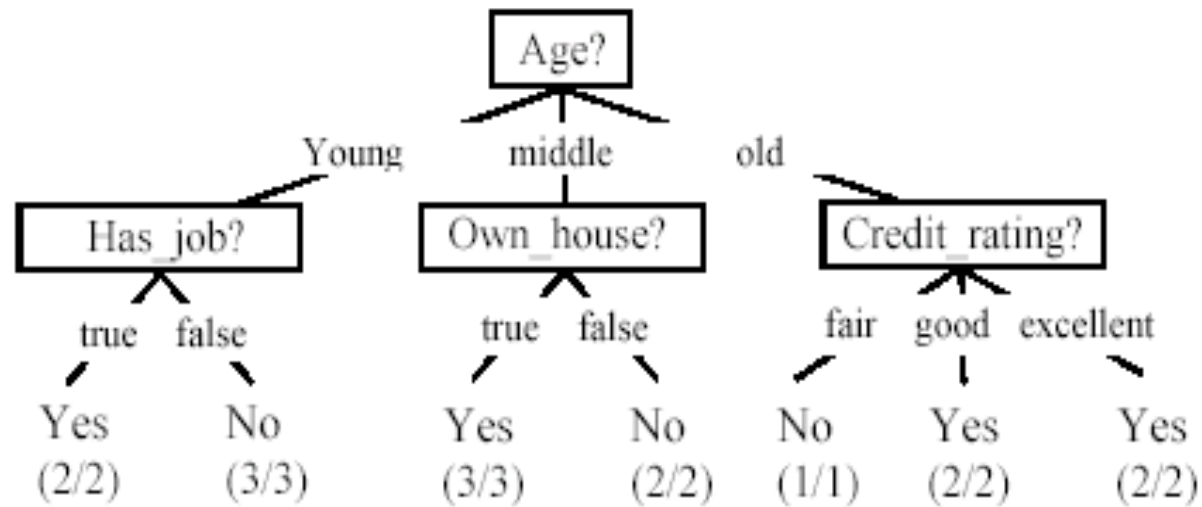


A DECISION TREE FOR THIS DATA

Decision nodes and leaf nodes (classes)



MULTIPLE VALID TREES ARE POSSIBLE



HANDLING CONTINUOUS ATTRIBUTES

Split into two (or more) intervals.


How to find the best threshold to divide?

Use **information gain** or **gain ratio** again


Sort the **values** in increasing order $\{v_1, v_2, \dots, v_r\}$,

Consider **possible thresholds** between adjacent values v_i and v_{i+1} .


Test possible thresholds and choose one that **maximizes the gain**.



AN EXAMPLE IN A CONTINUOUS SPACE



(A) A partition of the data space



(B). The decision tree

MANY FORMS OF DECISION TREES

By **Ross Quinlan**
– Most similar to
what we just saw

ID3 – Basic entropy-based decision trees (discrete only)

C4.5 – Handles continuous and discrete attributes

C5.0 – Enhanced version of C4.5

sk-learn
supports

CART (Classification And Regression Trees)

– Similar to C4.5 but uses “Gini Index” for purity instead of entropy

```
from sklearn.tree import DecisionTreeClassifier
```

And others ... CHAID/MARS/etc.

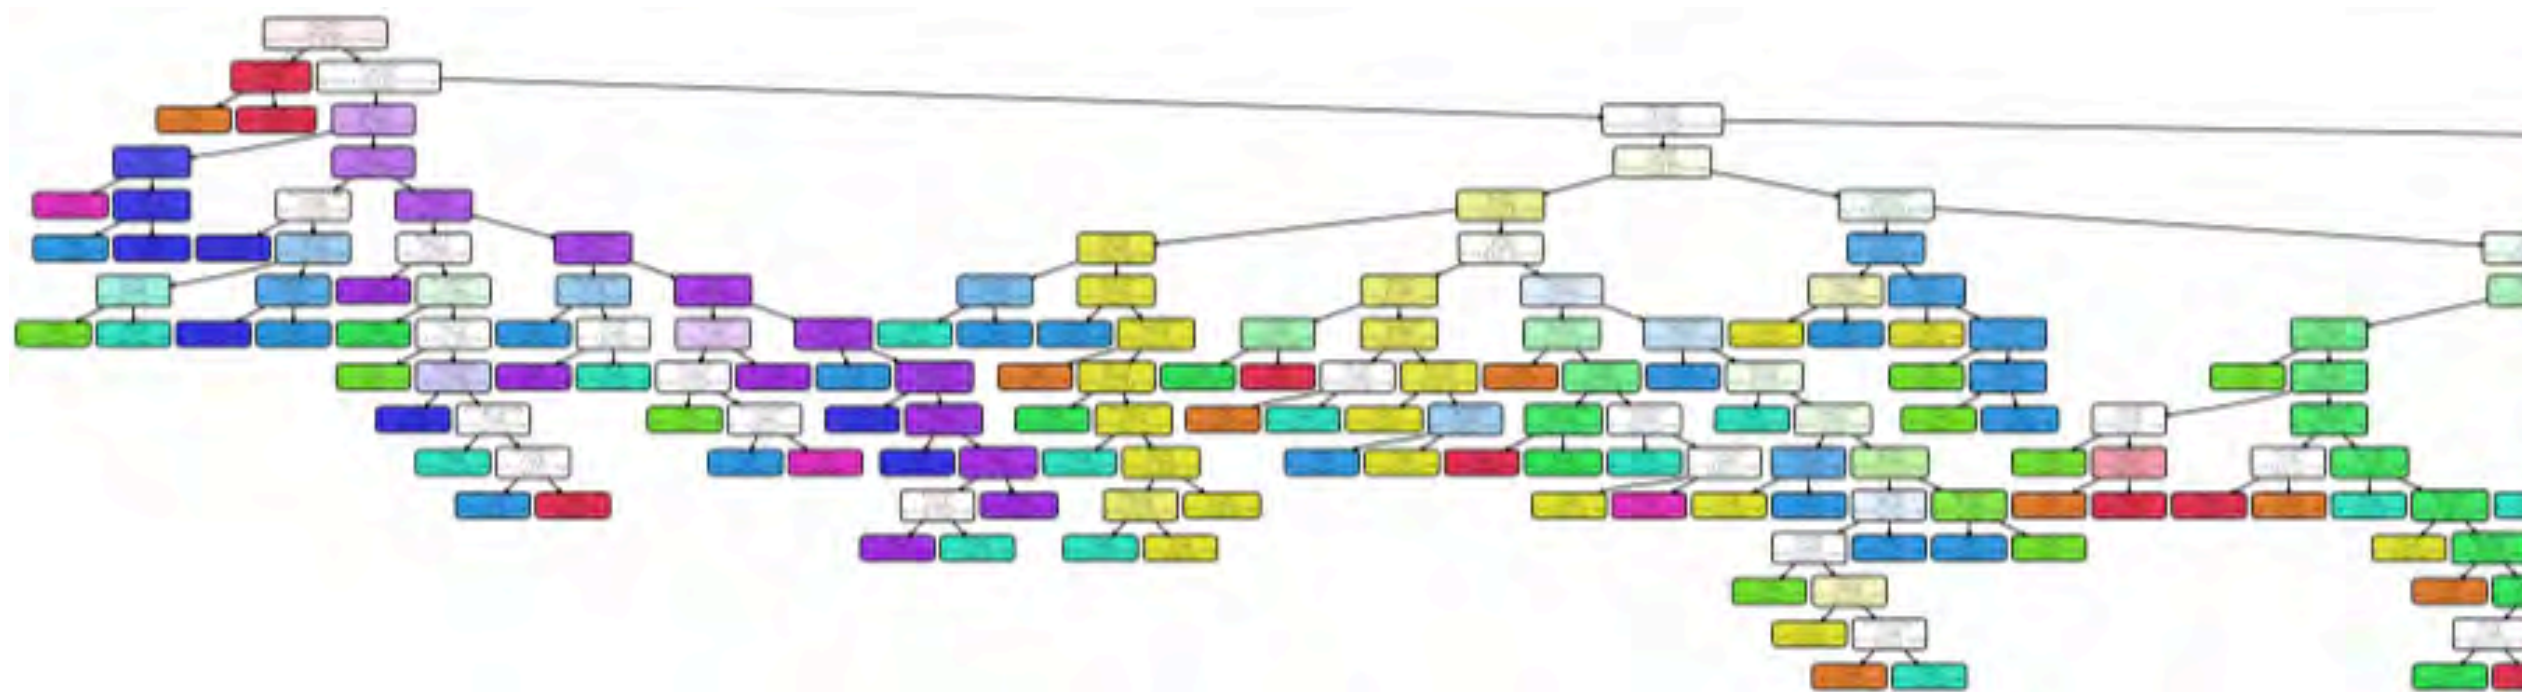
DECISION TREE

Advantages:


- Easy to interpret – not all classifiers can be “explained”
- Prediction process obvious
- Handle mixed data types

Disadvantages:



- Expensive to calculate
- Tendency to overfit
- Can get large



RISK OF OVERFITTING



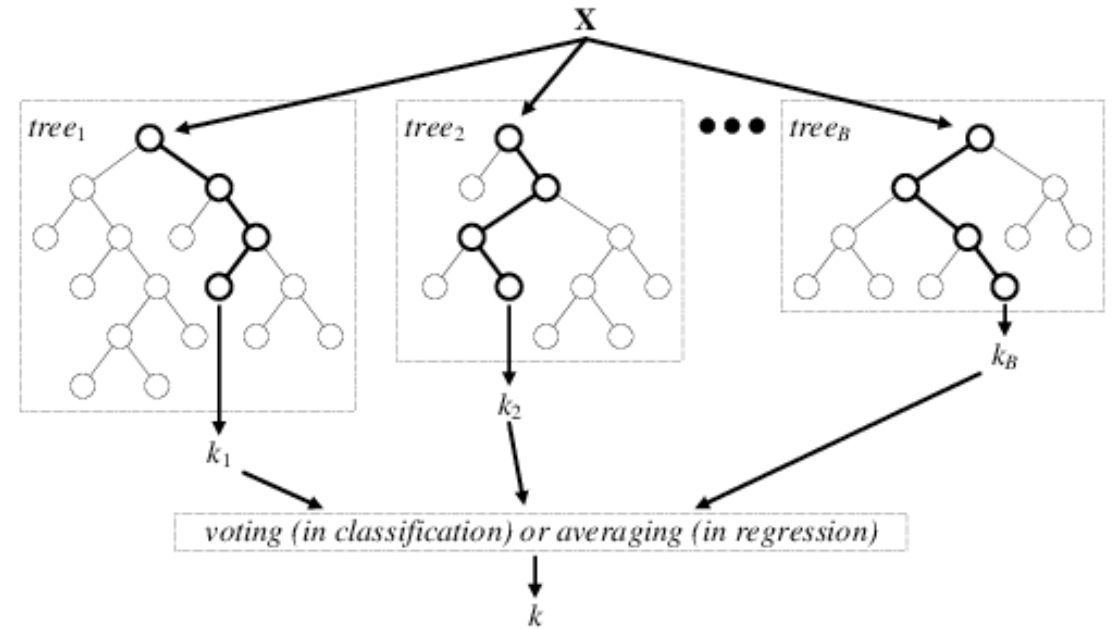
Different training samples can give really different results.



RANDOM FOREST

“Ensemble” classifiers

Create many trees and have them vote.



Verikas et al. (2016)

Problem: how to generate many trees from one dataset?

Various way of randomizing


Pick different data subsets (often using bootstrapping “with replacement”)

Pick different features


```
from sklearn.ensemble import RandomForestClassifier
```



```
from sklearn.ensemble import RandomForestClassifier  
model = RandomForestClassifier(n_estimators=100, random_state=0)
```



Single Decision Tree




Random Forest with 100 Trees

NEURAL NETS

(and “Deep Learning”)


MIMICING LEARNING IN THE BRAIN

“Neurons”



EACH NEURON OUTPUTS A WEIGHTED SUM OF ALL OUTPUTS FROM THE PREVIOUS LAYER

Multiple layers with weighted
“feed-forward” connections



Learning via back-propagation

EVERY EDGE HAS A UNIQUE WEIGHT THAT CAN CHANGE OVER TIME




Learning

A neural network getting to know faces (above) trains itself on perhaps millions of examples before it can pick out an individual face from a crowd or a cluttered landscape.

Input layer

Hidden layers

Output



Joe

Chris

Lee

Recognition

Input of a face into the network is analyzed at each layer before the network guesses correctly about its identity.



Each layer identifies progressively more complex features



LEARNING PREDICTABLY

Smooth activation functions

Sigmoid, Tanh, etc. (not stepwise)

Randomize weights initially

Adjust weights slowly





Epoch
000,000

Learning rate
0.03

Activation
Tanh

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

- X_1
- X_2
- X_1^2
- X_2^2
- X_1X_2
- $\sin(X_1)$

+ - 2 HIDDEN LAYERS

+ -

4 neurons

+ -

2 neurons


This is the output from one **neuron**.
Hover to see it larger.

The outputs are mixed with varying **weights**, shown by the thickness of the lines.


OUTPUT

Test loss 0.500

Training loss 0.503



Colors shows data, neuron and weight values.



☐ Show test data

☐ Discretize output

PLAYGROUND.TENSORFLOW.ORG

NEURAL NETS IN PYTHON

Supported in SciKitLearn (but not very scalable)

```
from sklearn.neural_network import MLPClassifier
```

A bunch of other libraries...

tensorflow

pylearn2

sklearn_theano

nolearn

keras

lasagne

...



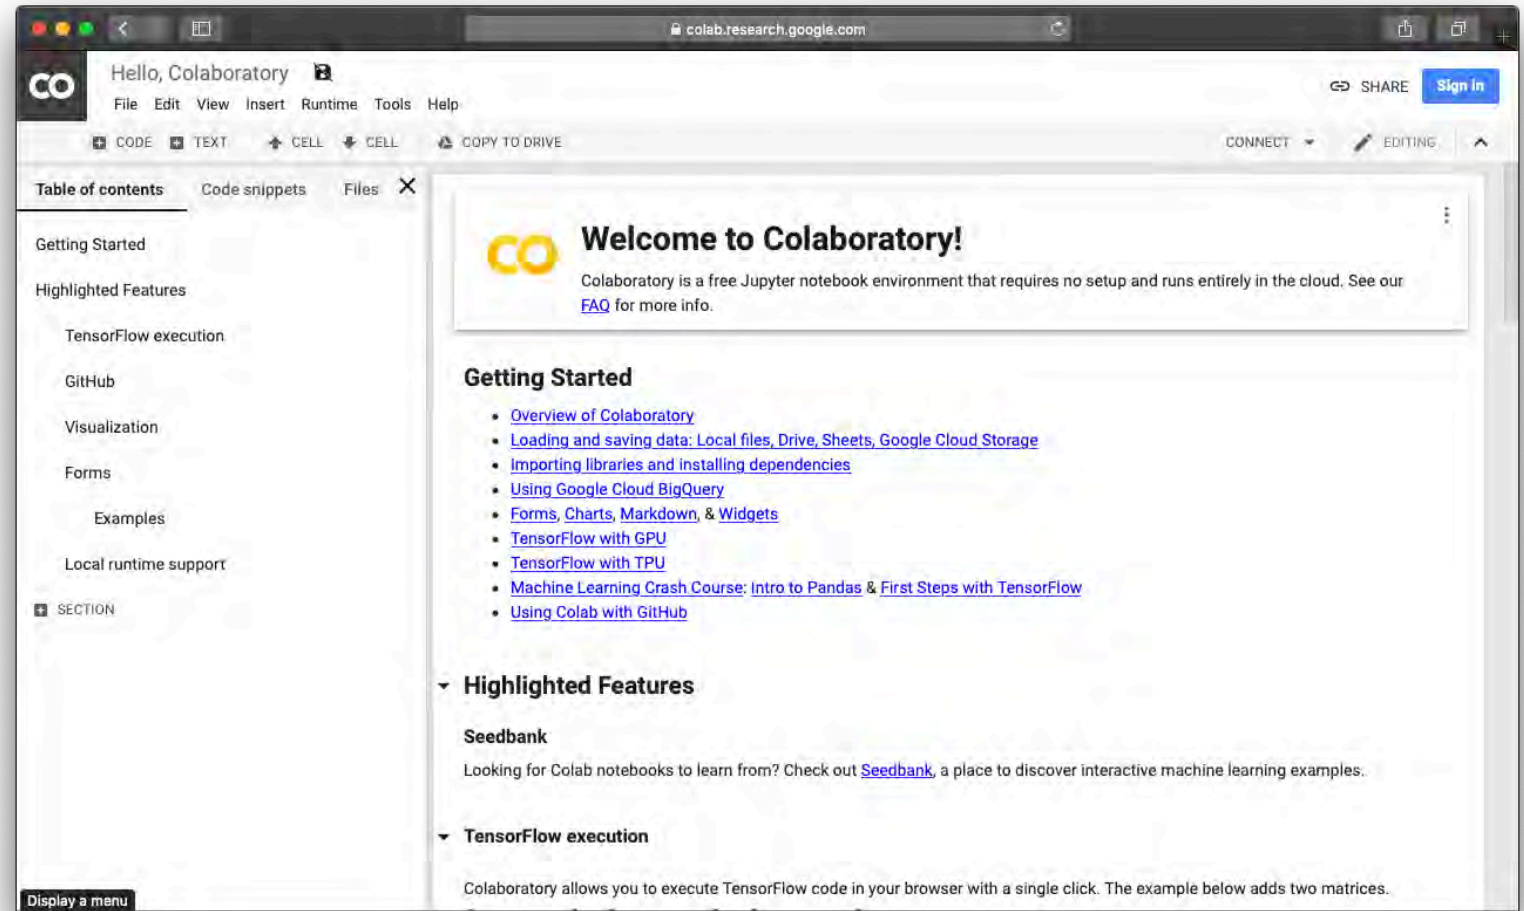
TensorFlow

TensorFlow is a fast and flexible open-source ML library for research and production.

colab

TensorFlow + Colaboratory

If you're going to start really diving into ML – this is a good place to start.



BUT... NEURAL NETS ARE BASICALLY BLACK BOXES

To perform well, NNs usually require
large numbers of nodes, multiple hidden layers, and lots of edges!

Internal behavior can be really difficult to understand.

Especially true with “deep learning”
(5, 10, or even more hidden layers).

OTHER APPROACHES

Logistic Regression

- Multiple linear regression extended to support categorical outputs instead of just quantitative ones.

```
from sklearn.linear_model import LogisticRegression
```

Support Vector Machines (SVM)

- Find a hyperplane in multidimensional space that best splits items with a label from items without it.

```
from sklearn.linear_model import LogisticRegression
```

ROLL-YOUR-OWN ENSEMBLE METHOD

Use a `VotingClassifier` to combine multiple classifiers.

```
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
```

```
clf1 = LogisticRegression()
clf2 = RandomForestClassifier()
clf3 = GaussianNB()
eclf = VotingClassifier(estimators=[('lr', clf1), ('rf', clf2), ('gnb', clf3)])
```



Can also weight models, tweak voting schemes, etc.

SUMMARY

Classification

Useful when we know something about the structure

Use a few labeled examples to classify many more

If you're interested – start playing
or take a machine learning course (like DATA 607)!

A NICE STARTING POINT

A nice, recent (Sept 2019)
Machine Learning reference built
around Python examples.

Available online via the UofC.

<https://learning.oreilly.com/library/view/hands-on-machine-learning/9781492032632/>

